



CS4379: Parallel and Concurrent Programming

CS5379: Parallel Processing

Lecture 2

Dr. Yong Chen

Associate Professor

Computer Science Department

Texas Tech University



Course Info

- **Lecture Time:** TR, 12:30-1:50
- **Lecture Location:** ECE 217
- **Sessions:** CS4379-001, CS4379-002, CS5379-001, CS5379-D01
- **Instructor:** Yong Chen, Ph.D., Associate Professor
- **Email:** yong.chen@ttu.edu
- **Phone:** 806-834-0284
- **Office:** Engineering Center 315
- **Office Hours:** 2-4 p.m. on Wed., or by appointment
- **TA:** TBA
- **More info:**
 - <http://www.myweb.ttu.edu/yonchen>
 - <http://discl.cs.ttu.edu>; <http://cac.ttu.edu/>; <http://nsfcac.org>



Outline

- Questions?
- Parallel architectures: classification based on control mechanism
- Parallel architectures: classification based on address space



Architectures

- Classification **based on control mechanism** (Flynn's Taxonomy)
 - SISD (Single Instruction Single Data stream)
 - SIMD (Single Instruction Multiple Data stream)
 - MISD (Multiple Instruction Single Data stream)
 - MIMD (Multiple Instruction Multiple Data stream)

- Classification **based on address space organization** (from programmers' point of view)
 - Shared Address Space
 - Distributed Address Space



Flynn's Taxonomy

- Flynn's taxonomy distinguishes parallel computer architectures according to how they can be classified along two independent dimensions of **Instruction** and **Data**
- Each of these dimensions can have only one of two possible states: **Single** or **Multiple**.
- The matrix below defines the 4 possible classifications according to Flynn:


SISD Single Instruction, Single Data	SIMD Single Instruction, Multiple Data
MISD Multiple Instruction, Single Data	MIMD Multiple Instruction, Multiple Data

Courtesy: Blaise Barney, LLNL



Single Instruction, Single Data (SISD)

- Model of serial Von Neumann machine (non-parallel)
- **Single instruction**: only one instruction stream is being executed by the CPU during any one clock cycle (single control processor)
- **Single data**: only one data stream is being used as input during any one clock cycle
- This is the oldest and even today, the most common type of computer till multicore era
- Examples: older generation mainframes, minicomputers and workstations



Dell laptop (uni-core)

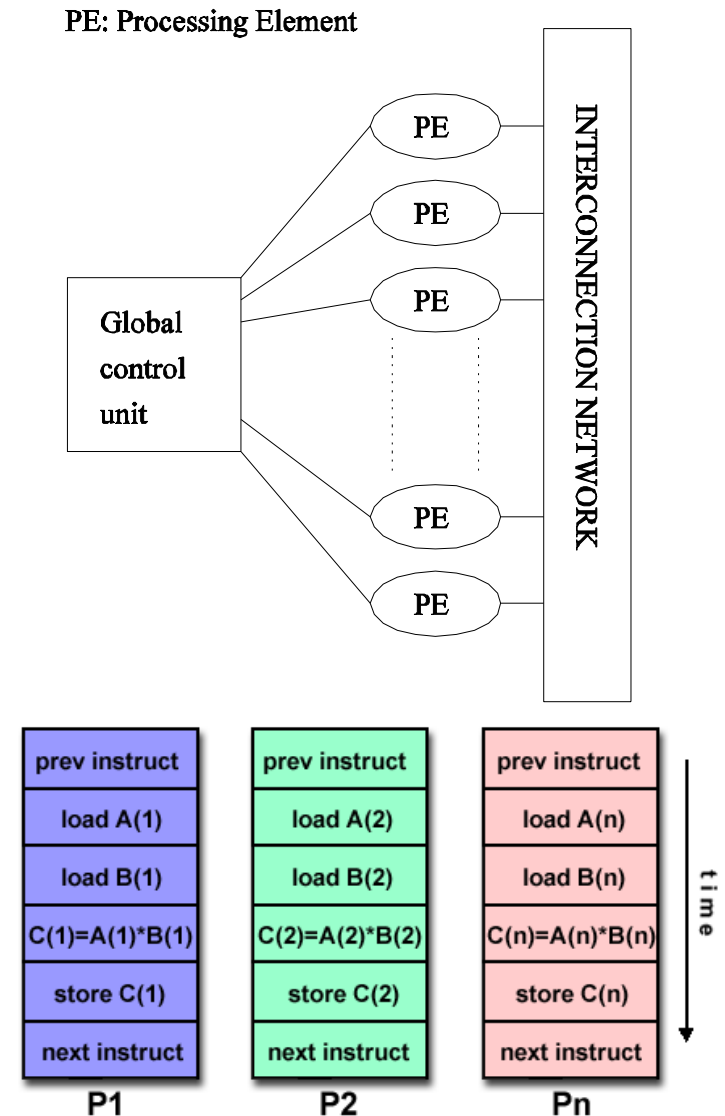


UNIVAC1




Single Instruction, Multiple Data (SIMD)

- **Single instruction**: All processing units execute the same instruction at any given clock cycle
- **Multiple data**: Each processing unit operate on a different data element
- It is often necessary to selectively turn off operations on certain data items. For this reason, most SIMD programming paradigms allow for an “**activity mask**”, which determines if a processor should participate in a computation or not
 - See the example in ITPC, Chapter 2.3



SIMD (cont.)

- Best suited for specialized problems with **high degree of regularity, such as graphics/image processing**
- Examples:
 - Connection Machine CM-2
 - IBM 9000, Cray X-MP, Y-MP & C90
- Most modern computers, **particularly GPUs employ SIMD instructions and execution units**



CM-2



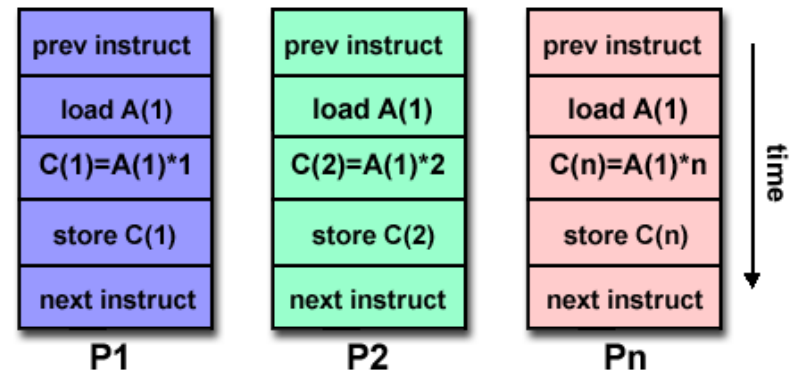
Cray X-MP

Fermi class
GTX480



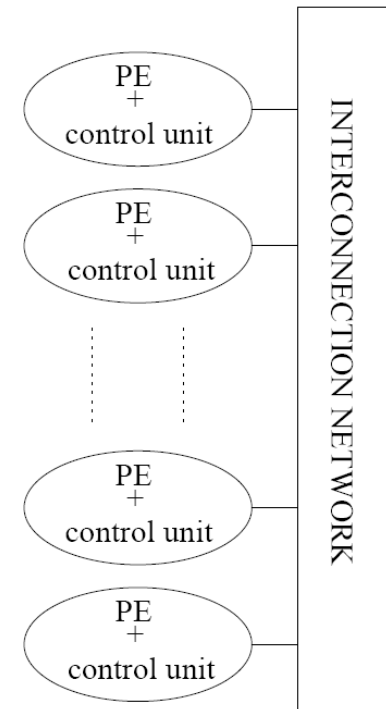
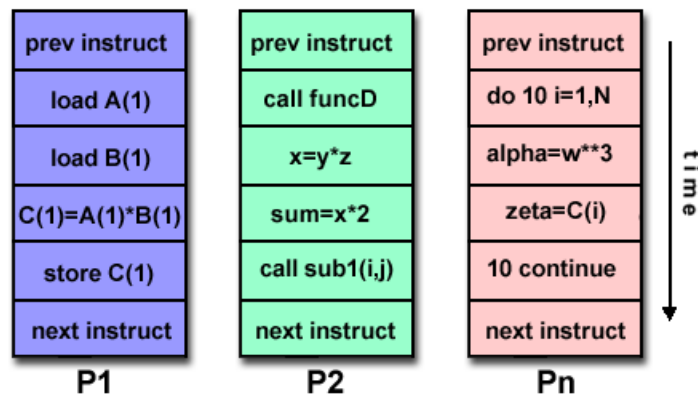
Multiple Instruction, Single Data (MISD)

- A single data stream is fed into multiple processing units
- Each processing unit operates on the data independently via independent instruction streams
- Few actual examples of this class of parallel computer have ever existed
- Some conceivable uses might be
 - Multiple frequency filters operating on a single signal stream
 - Multiple cryptography algms attempting to crack a single coded msg



Multiple Instruction, Multiple Data (MIMD)

- Most common type of parallel computer
- **Multiple Instruction**: every processor may execute different instruction stream
 - May have separate clocks
- **Multiple Data**: every processor may be working with a different data stream





MIMD (cont.)

- Examples
 - ❑ IBM SP, TMC's CM-5, Cray T3D & T3E, SGI Origin, Tera MTA, ...
 - ❑ Most current supercomputers
 - ❑ Networked parallel computer clusters
 - ❑ Multi-processor SMP (symmetric multiprocessing) computers
 - ❑ Multi-core PCs
- Note: many MIMD architectures also include SIMD execution sub-components (e.g. GPUs)





Single Program, Multiple Data (SPMD)

- A variant of MIMD, executes the same program (multiple instances) on different processors
 - MIMD processors can execute different programs on different processors
- Widely used by many parallel platforms
 - E.g. executions of MPI (message passing interface) programs



SIMD v.s. MIMD

■ SIMD: Pros v.s. Cons?

- ❑ SIMD computers **require less hardware** than MIMD computers (single control unit), **less memory** (only one copy of the program)
- ❑ However, since SIMD processors are specially designed, they tend to be expensive and have long design cycles
- ❑ **Not all applications are naturally suited to SIMD processors**

■ MIMD: Pros v.s. Cons?

- ❑ In contrast, MIMD computers can be built from inexpensive off-the-shelf components with relatively little effort in a short amount of time
- ❑ **Suitable for the irregular nature of many applications**
- ❑ Store the program and OS at each processor, **requires more hardware**




Outline

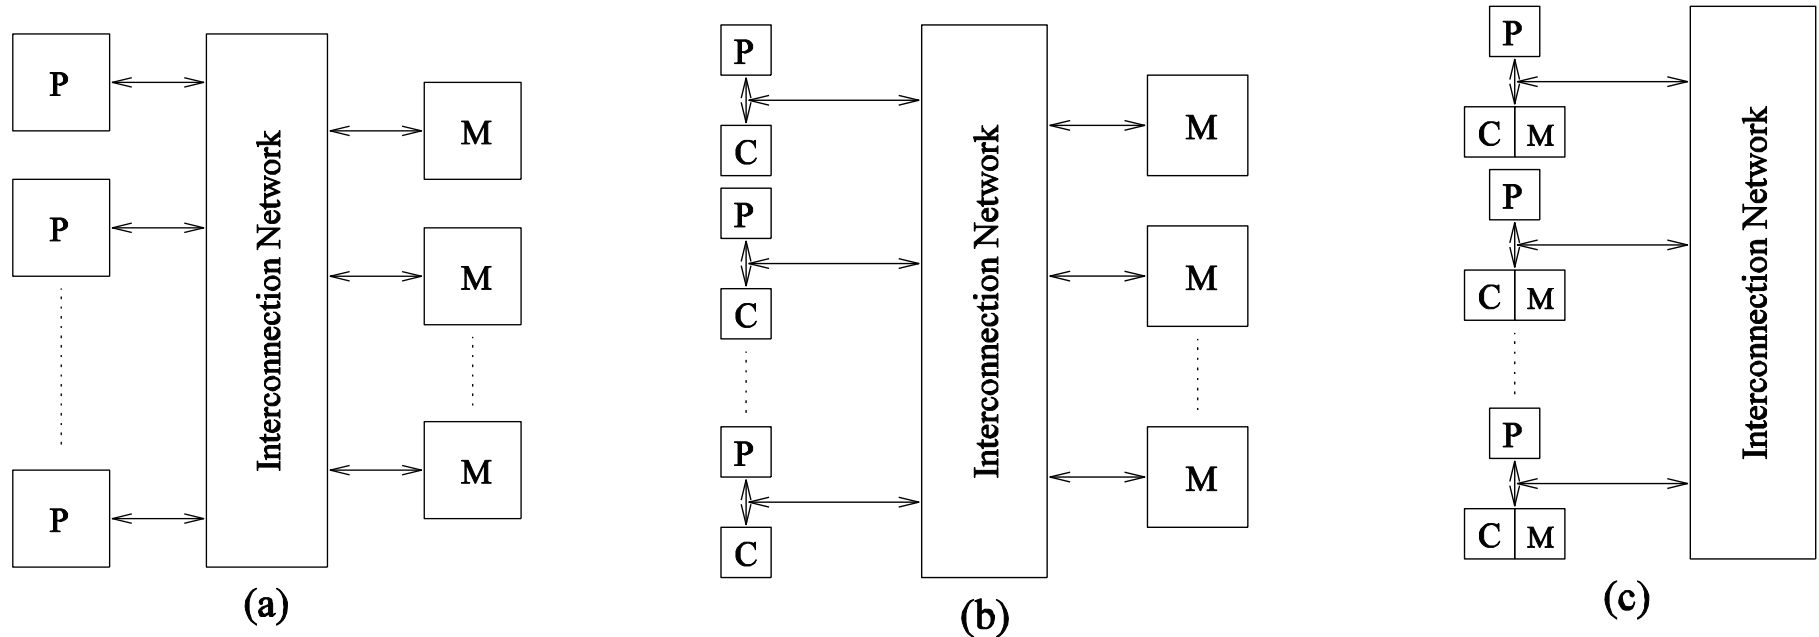
- Questions?
- Parallel architectures: classification based on control mechanism
- Parallel architectures: classification based on address space

Shared-Address-Space Architecture

- Processors (or processing units, or processing elements) can directly access all the data in the system
- Interaction?
 - Processors interact by modifying data stored in this shared-address-space
- If the time taken by a processor to access any memory word in the system global or local is identical, the platform is classified as a **uniform memory access (UMA)**, else, a **non-uniform memory access (NUMA)** machine.



NUMA and UMA Shared-Address-Space Architectures



Typical shared-address-space architectures:

- (a) **UMA** shared-address-space computer;
- (b) **UMA** shared-address-space computer **with caches and memories**;
- (c) **NUMA** shared-address-space computer with local memory only.



NUMA and UMA

Shared-Address-Space Architectures

- The distinction between NUMA and UMA platforms is important from the point of view of algorithm design. NUMA machines require locality from underlying algorithms for performance.
- Read/write shared data must be coordinated (this will be discussed in detail when we talk about threads programming).
- Caches in such machines require coordinated access to multiple copies. This leads to the cache coherence problem.
- A weaker model of these machines provides an address map, but not coordinated access. These models are called non cache coherent shared address space machines.
 - ccNUMA v.s. nccNUMA

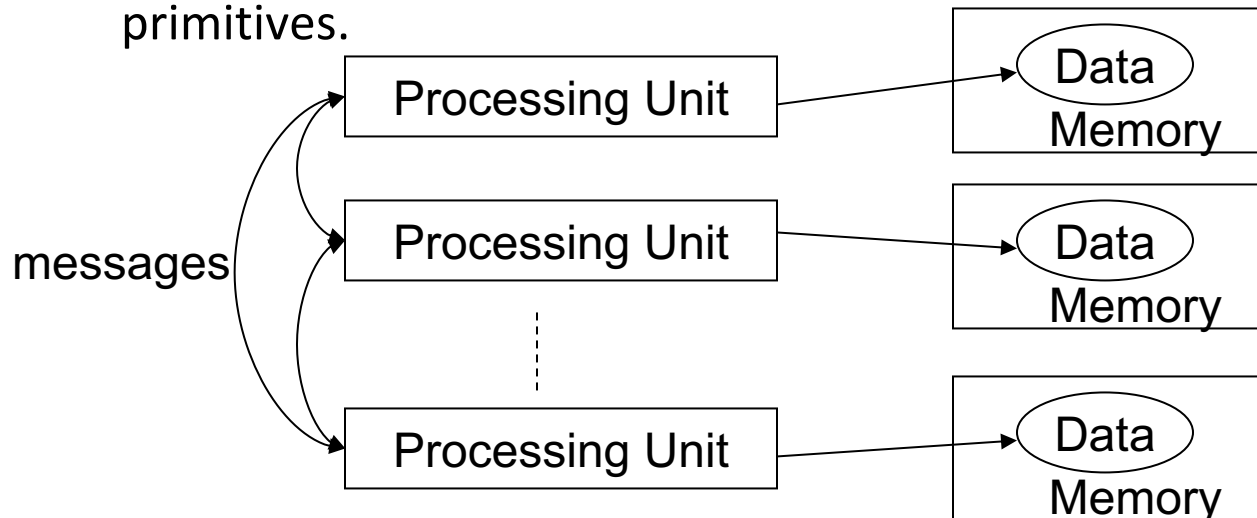


Shared-Address-Space vs. Shared Memory Machines

- It is important to **note the difference** between the terms **shared address space** and **shared memory**.
- We refer to the former as a **programming abstraction** and to the latter as a **physical machine attribute**.
- It is possible to provide a shared address space using a physically distributed memory.
 - **Distributed Shared Memory (DSM) architecture**
 - Identical to NUMA


Distributed-Address-Space Architectures

- “Shared nothing:” each processor has a private address space
- Processors can directly access only local data
 - ❑ No global address space
- Interaction?
 - ❑ These platforms are programmed using (variants of) message passing primitives, e.g. send and receive
 - ❑ Libraries such as MPI, sockets, MapReduce provide such primitives.






Distributed Address Space



P: Processor
M: Memory



A Hrothgar@TTU Parallel Computer Example





Shared vs. Distributed Address Space

- Shared address space: Pro. Vs Con.?
 - Global address space view for programmers, easier to program usually
 - Implicit communication
 - Need hardware/software to support view, complexity in system design
 - Not easy to be scalable

- Distributed address space: Pro. Vs. Con.?
 - Requires little hardware support, other than a network
 - Easy to scale up
 - No global address space view, more difficult to program
 - Need explicit communication

- Shared address space platforms can easily emulate message passing. The reverse is more difficult to do (in an efficient manner).



Multiprocessors v.s. Multicomputers

- “Multiprocessors”: platforms that provide shared-address-space
- “Multicomputers”: platforms that do not provide shared-address space and need message passing communications



- 24



Questions?

Questions/Suggestions/Comments are always welcome!

Write me: yong.chen@ttu.edu

Call me: 806-834-0284

See me: ENGCTR 315

If you write me an email for this class, please start the email subject with [CS4379] or [CS5379].