



# **CS4379: Parallel and Concurrent Programming**

## **CS5379: Parallel Processing**

### **Lecture 7**

**Dr. Yong Chen**

**Associate Professor**

**Computer Science Department**

**Texas Tech University**



## Course Info

- **Lecture Time:** TR, 12:30-1:50
- **Lecture Location:** ECE 217
- **Sessions:** CS4379-001, CS4379-002, CS5379-001, CS5379-D01
- **Instructor:** Yong Chen, Ph.D., Associate Professor
- **Email:** [yong.chen@ttu.edu](mailto:yong.chen@ttu.edu)
- **Phone:** 806-834-0284
- **Office:** Engineering Center 315
- **Office Hours:** 2-4 p.m. on Wed., or by appointment
- **TA:** Mr. Ghazanfar Ali, [Ghazanfar.Ali@ttu.edu](mailto:Ghazanfar.Ali@ttu.edu)
- **TA Office hours:** Tue. and Fri., 2-3 p.m., or by appointment
- **TA Office:** EC 201 A
- **More info:**
  - <http://www.myweb.ttu.edu/yonchen>
  - <http://discl.cs.ttu.edu>; <http://cac.ttu.edu/>; <http://nsfcac.org>



## Announcements

- Two guest lectures by Mr. Misha Ahmadian next week on the subject of how to use parallel computers on campus
- If you have a laptop please consider bringing it with you and you can follow the lecture to try to use the systems



# Outline

- Questions?
- Amdahl's law revisited and scaled speedup
- Scalability of parallel systems
- Performance evaluation and analysis tools
- Quiz #2



## Amdahl's Law

- Tacit assumption in Amdahl's law
  - The problem size or workload,  $W$ , is **fixed**
  - The speedup emphasizes **time reduction**
  
- **Fixed-Size Speedup**

$$S_p = \frac{\text{Uniprocessor Execution Time}}{\text{Parallel Execution Time}}$$

$$S_p = \frac{\text{Uniprocessor Time of Solving } W}{\text{Parallel Time of Solving } W}$$



## Amdahl's Law



- Amdahl's analysis argues a limit on speedup in terms of  $\alpha$

$$T_p = \alpha T_s + \frac{(1 - \alpha)T_s}{p}$$

$$S_p = \frac{T_s}{\alpha T_s + \frac{(1 - \alpha)T_s}{p}} = \frac{1}{\alpha + \frac{1 - \alpha}{p}}$$

$$\lim_{p \rightarrow \infty} S_p = \frac{1}{\alpha}$$

# Fixed-Size Speedup





# Impact of Amdahl's Law



IBM 7030 Stretch



IBM 7950 Harvest

All have up to 8/16/32 processors, citing Amdahl's law,

$$\lim_{p \rightarrow \infty} \text{Speedup}_{\text{Amdahl}} = \frac{1}{\alpha}$$


Cray X-MP  
Fastest computer 1983-1985



Cray Y-MP






## Scaled Speedup Model

- **Fixed-Time Speedup (Scaled Speedup)** (Gustafson, 88)
  - Emphasis on work finished in a fixed time
  - Problem size is scaled from  $W$  to  $W'$
  - $W'$ : Work finished within the fixed time with parallel processing

$$\begin{aligned} S'_p &= \frac{\text{Uniprocessor Time of Solving } W'}{\text{Parallel Time of Solving } W'} \\ &= \frac{\text{Uniprocessor Time of Solving } W'}{\text{Uniprocessor Time of Solving } W} \\ &= \frac{W'}{W} \end{aligned}$$



# Gustafson's Law

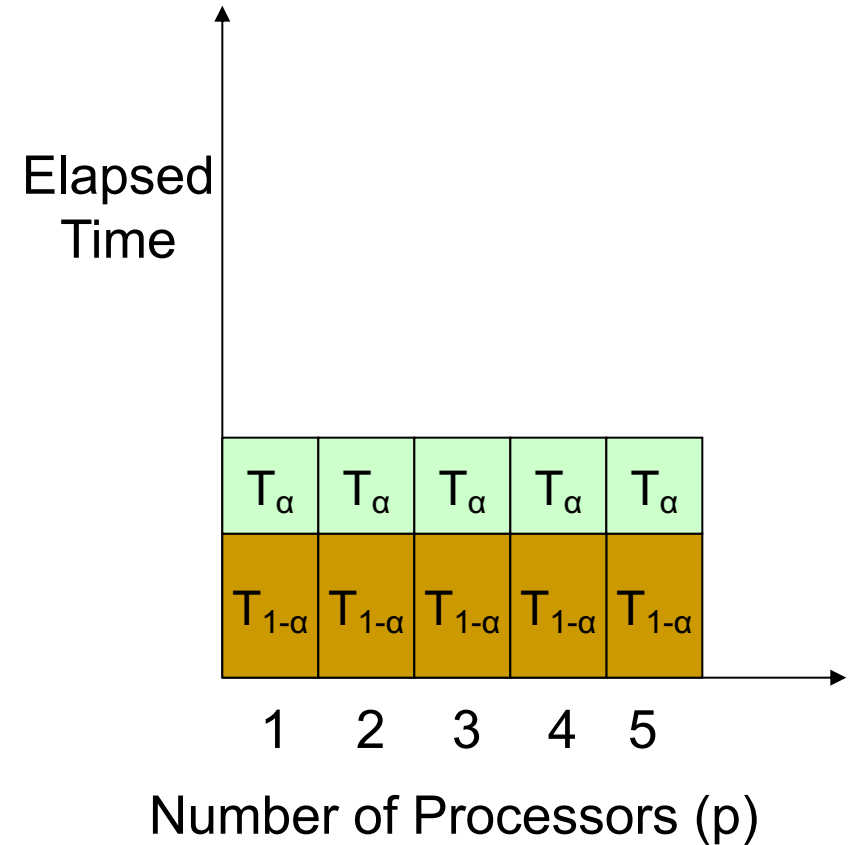



$$Speedup_{FT} = \frac{Work(p)}{Work(1)} = \frac{\alpha W + (1-\alpha)pW}{W} = \alpha + (1-\alpha)p$$

*(FT stands for fixed-time)*




## Fixed-Time Speedup (Gustafson's Law)

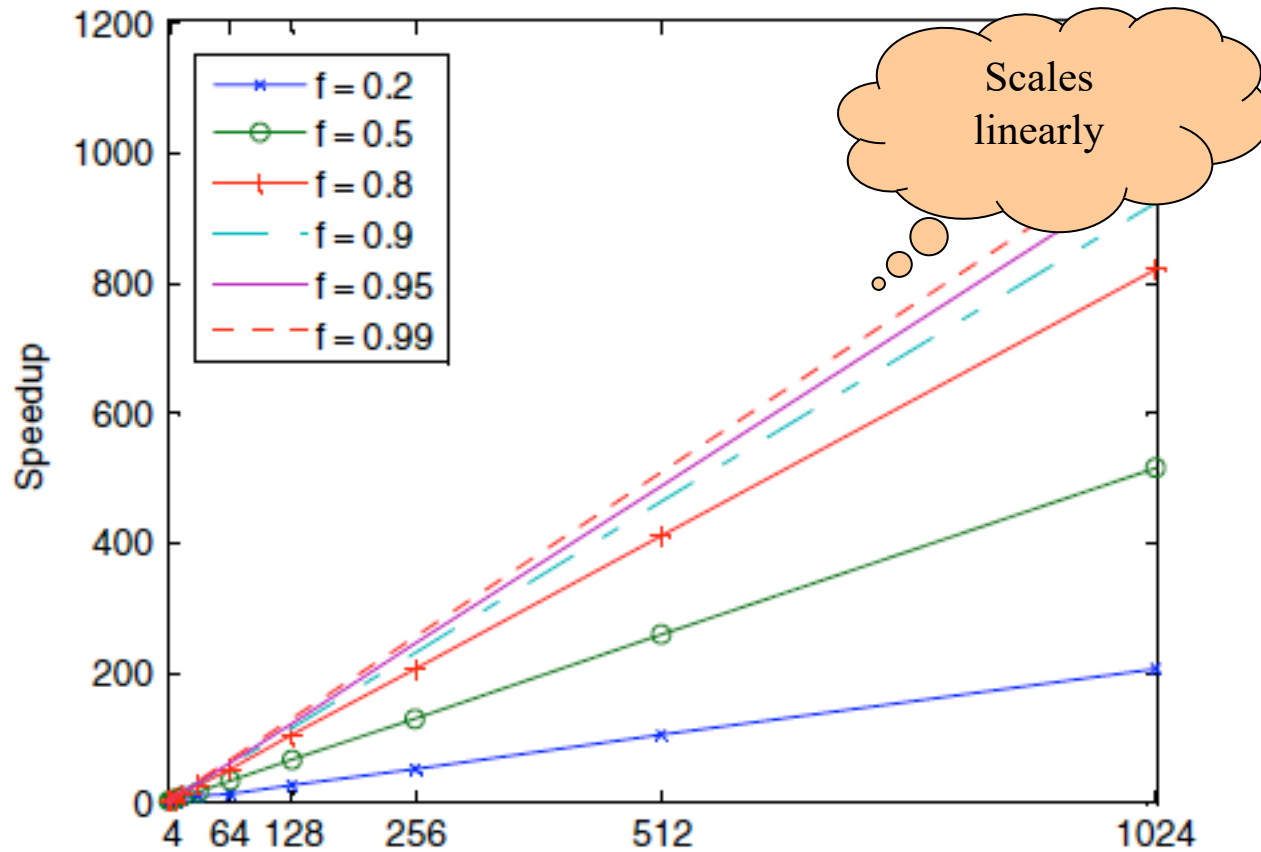




## Fixed-size Speedup



# Fixed-time Speedup





# Terminologies

- **Strong scaling**: problem size fixed (i.e. what Amdahl's law focuses)
- **Weak scaling**: problem size scaled (i.e. what Gustafson's law focuses)
  - E.g. fixed-time speedup model
  - Other scaled speedup models also exist, e.g. memory-bounded scaled speedup model



# Why Scaled Computing

- Solve larger problems
  - May not fit or can't run for small machine
- Better solution, better accuracy
  - e.g. computing pi
- Maintain efficiency
- Provide real-time solution
  - May not be achievable with small machines



# Scaled Computing Leads to Today's Systems



LANL Roadrunner:  
25,200 processors, 2008  
World's fastest supercomputer

The system scale is  
**far beyond**  
implication of  
Amdahl's law

TACC Ranger:  
15,744 processors,  
2008



ANL Intrepid:  
20,480 processors, 2008

Note: data years ago





# Outline

- Questions?
- Amdahl's law revisited and scaled speedup
- Scalability of parallel systems
- Performance evaluation and analysis tools



# Scaling Characteristics of Parallel Programs

- The efficiency of a parallel program can be written as:

$$E = \frac{S}{p} = \frac{T_S}{pT_P}$$

$$E = \frac{1}{1 + \frac{T_o}{T_S}}$$

or

$$T_o = p T_P - T_S$$



# Scaling Characteristics of Parallel Programs

- For a given problem size (i.e., the value of  $T_s$  remains constant), as we increase  $p$  (i.e. the number of processing elements, or the **system size**),  $T_o$  increases
  - **Problem size (workload)  $W$** : total number of operations required to solve a problem, e.g.  $2N^3$  for matrix multiplication
  - i.e. **the total overhead function  $T_o$  is an increasing function of  $p$**
- Thus the overall efficiency of the parallel program goes down

*Matrix\_multiplication(A, B, C)*

```
for (i = 0; i < n; i++)  
    for (j = 0; j < n; j++)  
        for (k = 0; k < n; k++)  
            C[i, j] = C[i, j] + A[i, k] * B[k, j]
```


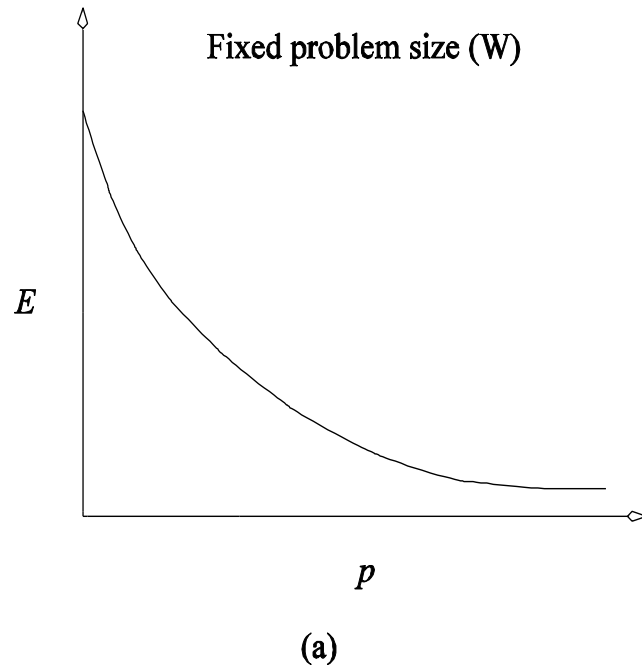


# Scaling Characteristics of Parallel Programs

- The total overhead function  $T_o$  is also a function of problem size  $W$
- In many cases,  $T_o$  grows sub-linearly with respect to  $W$
- In such cases, the efficiency increases if the problem size is increased, given the system size (the number of processing elements) remains constant



## Variation of efficiency



Variation of efficiency: (a) as the number of processing elements is increased for a given problem size; and (b) as the problem size is increased for a given number of processing elements.



# Scalable Parallel Systems

- We can simultaneously increase the problem size and the number of processors to keep efficiency constant.
- **Scalable** parallel systems: with the ability to maintain the efficiency at a fixed value by simultaneously increasing the number of processing elements (system size) and problem size



## Isoefficiency Metric of Scalability

- What is the rate at which the problem size must increase with respect to the number of processing elements to keep the efficiency fixed?
- This rate determines the scalability of the system. The slower this rate, the better.
- **Scalability** of parallel systems: a measure of its capacity to increase speedup (in order to keep efficiency constant) in proportion to the number of processors given increasing the problem size
  - Isoefficiency scalability (see details from additional paper readings)



# Scalability of Heterogeneous Computing

- What is a heterogeneous parallel system?
- Is the previously discussed scalability metric well applicable to heterogeneous computing?
- No, as system scaling cannot be well represented by  $p$  anymore
  - The efficiency definition is an issue
- System scaling: marked speed
- Contains homogeneous as a special case

Y. Chen, X.-H. Sun and M. Wu. Algorithm-System Scalability of Heterogeneous Computing. Journal of Parallel and Distributed Computing (JPDC), Vol. 68, No. 11, 1403 – 1412, 2008.





# Outline

- Questions?
- Amdahl's law revisited and scaled speedup
- Scalability of parallel systems
- Performance evaluation and analysis tools
- Quiz #2



# Performance Evaluation and Analysis

## ■ Performance measurement

- Define metrics (run time, speedup, scalability), adjust parameters / variables (controllable and non-controllable, e.g. noises, jitter, variation, etc.)

## ■ Performance modeling, prediction

- Build up a model, calculate/analyze the performance, like the  $\pi$  example we show earlier
- A model can be used to predict the performance on a platform we don't have

## ■ Performance diagnosis/optimization

- Instrumentation/profiling
- Post-execution, traces

Often combine these approaches too



# Measurement

- Timing an entire program
  - UNIX *time* command outputs
    - User time
    - System time
    - Elapsed time
  - User time + system time = CPU time
  - Additional *time* output
    - Percent utilization
    - Average memory utilization
    - Blocked I/O operations
    - Page faults and swaps



## Timing a Portion of a Program

- Record the time before you start doing  $x$
- Do  $x$
- Record the time at completion of  $x$
- Subtract the start time from the completion time
- Often repeat  $x$  with  $i$  iterations (e.g. a million iterations), then divide the time by  $i$



# Profiling

- Most compilers provide a facility to automatically insert timing calls into your program at the entry and exit of each routine at compile time
- A separate utility (e.g. *prof*, *gprof*) produces a report showing the percentage of time spent in each routine
- Many performance analysis tools also provide this capability



# Types of Profiling

- Time-based
- Based on other metrics such as
  - Operation counts
  - Cache and memory event counts

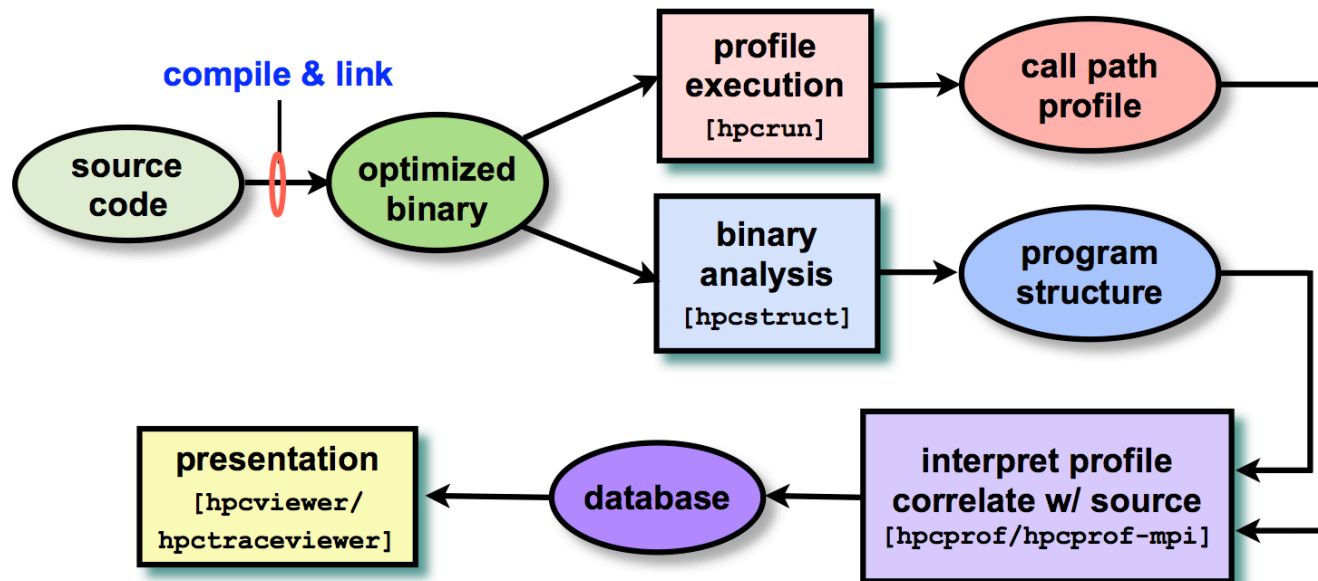


## PAPI

- <http://icl.cs.utk.edu/papi/>
- Provide support for accessing hardware performance counters available on most modern microprocessors

# HPCToolkit

- <http://hpctoolkit.org/>
- An open-source suite of multi-platform tools for profile-based performance analysis of applications.

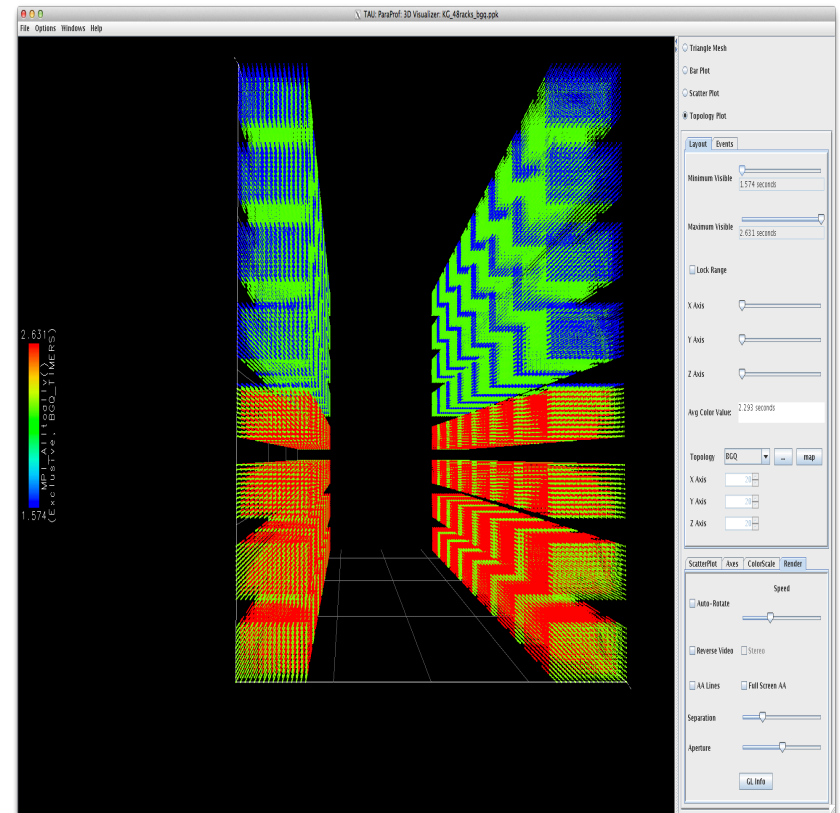






# TAU

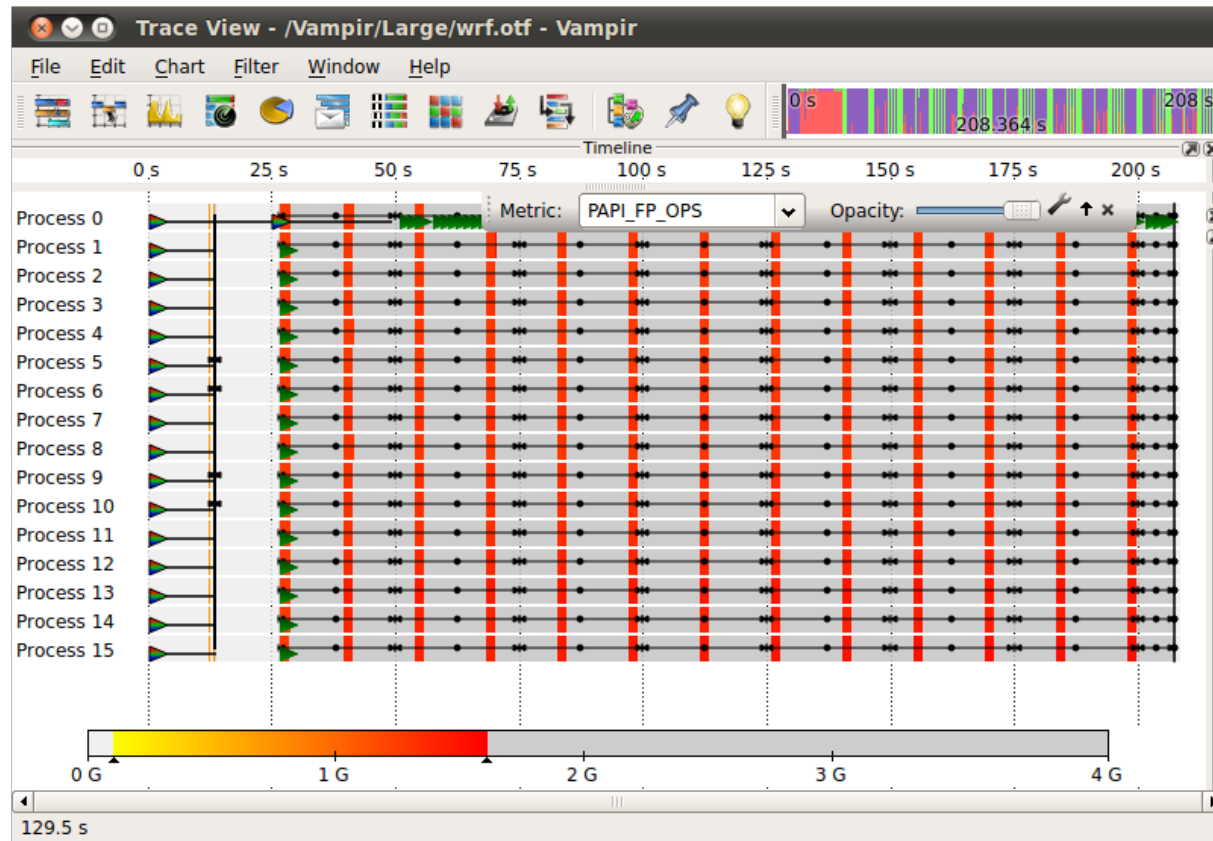
- <http://www.cs.uoregon.edu/research/tau/home.php>
- A portable profiling and tracing toolkit for performance analysis of parallel programs written in Fortran, C, C++, UPC, Java, Python





# Vampir

- <https://vampir.eu/>
- A performance visualization and optimization tool





## Performance Analysis Tools

- MPE logging.Jumpshot
- Pablo
- Paradyn
- Scalea
- VProf
- Prophecy
- DynaProf
- .....



## Readings

- Reference book ITPC, Chapter 5
  - 5.1, 5.2, 5.3 (more examples), 5.4, 5.6, 5.7, 5.8
- Foster, DBPP, Chapter 3, 3.1-3.4
  - <http://www.mcs.anl.gov/~itf/dbpp/text/book.html>
- John L. Gustafson, “*Reevaluating Amdahl’s Law*”, 1988
- A. Y. Grama, A. Gupta, and V. Kumar. Isoefficiency : Measuring the Scalability of Parallel Algorithms and Architectures, 1993.
- Y. Chen, X.-H. Sun and M. Wu. Algorithm-System Scalability of Heterogeneous Computing. Journal of Parallel and Distributed Computing (JPDC), Vol. 68, No. 11, 2008. (Extended ICPP-2005)



## Questions?

Questions/Suggestions/Comments are always welcome!

Write me: [yong.chen@ttu.edu](mailto:yong.chen@ttu.edu)

Call me: 806-834-0284

See me: ENGCTR 315

*If you write me an email for this class, please start the email subject with [CS4379] or [CS5379].*