



CS4379: Parallel and Concurrent Programming

CS5379: Parallel Processing

Lecture 15

Guest Lecture by Wei Zhang

X-Spirit.Zhang@ttu.edu

Dr. Yong Chen

Associate Professor

Computer Science Department

Texas Tech University




Course Info

- **Lecture Time:** TR, 12:30-1:50
- **Lecture Location:** ECE 217
- **Sessions:** CS4379-001, CS4379-002, CS5379-001, CS5379-D01
- **Instructor:** Yong Chen, Ph.D., Associate Professor
- **Email:** yong.chen@ttu.edu
- **Phone:** 806-834-0284
- **Office:** Engineering Center 315
- **Office Hours:** 2-4 p.m. on Wed., or by appointment
- **TA:** Mr. Ghazanfar Ali, Ghazanfar.Ali@ttu.edu
- **TA Office hours:** Tue. and Fri., 2-3 p.m., or by appointment
- **TA Office:** EC 201 A
- **More info:**
 - <http://www.myweb.ttu.edu/yonchen>
 - <http://discl.cs.ttu.edu>; <http://cac.ttu.edu/>; <http://nsfcac.org>



Pthreads Example Code: Compute pi

- Generating random points in a unit length square
- Counting the number of points fall within circle
- The fraction of random points: $\pi/4$





Pthreads Example Code: Compute pi

- Assigns a fixed number of points to each thread
- Each thread generates these random points
- Collecting the number of points land in the circle
- After all threads finish, all threads counts are combined to compute pi
- Calculate fraction over all threads and multiply by 4



Pthreads Example Code: Compute pi

```
#include <pthread.h>
#include <stdlib.h>
#define MAX_THREADS 512
void *compute_pi (void *);
....
main() {
    ...
    pthread_t p_threads[MAX_THREADS];
    pthread_attr_t attr;
    pthread_attr_init (&attr);
    for (i=0; i< num_threads; i++) {
        hits[i] = i;
        pthread_create(&p_threads[i], &attr, compute_pi,
            (void *) &hits[i]);
    }
    for (i=0; i< num_threads; i++) {
        pthread_join(p_threads[i], NULL);
        total_hits += hits[i];
    }
    ...
}
```



Pthreads Example Code: Compute pi

```
void *compute_pi (void *s) {
    int seed, i, *hit_pointer;
    double rand_no_x, rand_no_y;
    int local_hits;
    hit_pointer = (int *) s;
    seed = *hit_pointer;
    local_hits = 0;
    for (i = 0; i < sample_points_per_thread; i++) {
        rand_no_x =(double)(rand_r(&seed))/(double)((2<<14)-1);
        rand_no_y =(double)(rand_r(&seed))/(double)((2<<14)-1);
        if (rand_no_x * rand_no_x +
            rand_no_y * rand_no_y <= 1.0)
            local_hits ++;
        seed *= i;
    }
    *hit_pointer = local_hits;
    pthread_exit(0);
}
```



Programming Notes

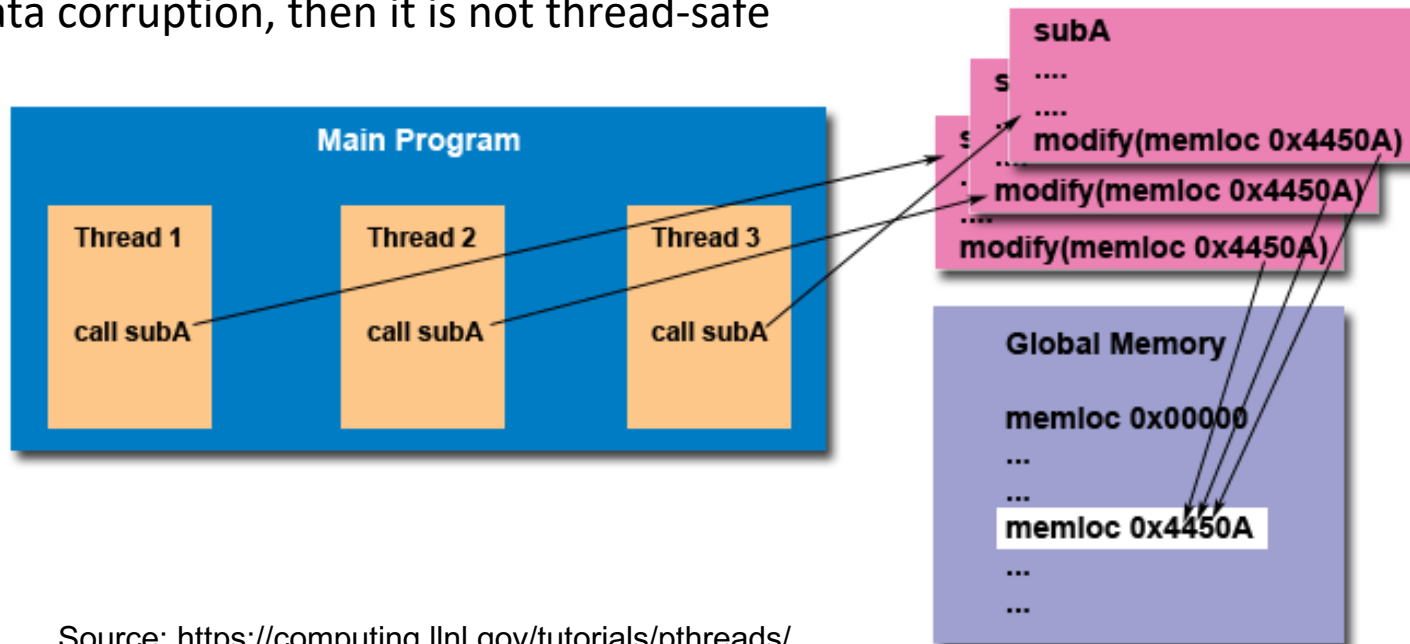
- Note the use of the function *rand_r* (instead of superior random number generators such as `drand48`)
- Reentrant functions: thread safety (thread-safe code)
- Can be safely called when another instance has been suspended in the middle of its invocation
- Thread can be preempted
- Non-reentrant function does not work as desired if another thread executes the same function
- Reading: [https://en.wikipedia.org/wiki/Reentrancy_\(computing\)](https://en.wikipedia.org/wiki/Reentrancy_(computing))



Programming Notes

■ Thread-safety:

- ❑ E.g. an app creates several threads, each making a call to the same library routine
- ❑ This library routine accesses/modifies a global structure or location in memory
- ❑ As each thread calls this routine it is possible that they may try to modify this global structure/memory location at the same time
- ❑ If the routine does not employ some sort of synchronization constructs to prevent data corruption, then it is not thread-safe



Source: <https://computing.llnl.gov/tutorials/pthreads/>



Outline

- Questions?
- Overview of programming models and thread basics
- The POSIX Threads
- Synchronization Primitives in Pthreads
 - Mutual Exclusion for Shared Variables
 - Producer-Consumer problem
 - Condition Variables for Synchronization
 - Controlling Thread and Synchronization Attributes
 - Composite Synchronization Constructs
- Compiling and running Pthreads program



Synchronization Primitives in Pthreads

- When multiple threads attempt to manipulate a shared data item, the results can often be incoherent if proper care is not taken to synchronize them
- Consider:

```
/* each thread tries to update a shared variable
   best_cost as follows */
if (my_cost < best_cost)
    best_cost = my_cost;
```
- Assume that there are two threads, the initial value of best_cost is 100, and the values of my_cost are 50 and 75 at threads t1 and t2.
- Depending on the schedule of the threads, the value of best_cost could be 50 or 75!



Synchronization Primitives in Pthreads

- Two problems
 - Parallel/concurrent nature (non-deterministic)
 - The possible value of 75 is **inconsistent** - it does not correspond to any serialization of the threads

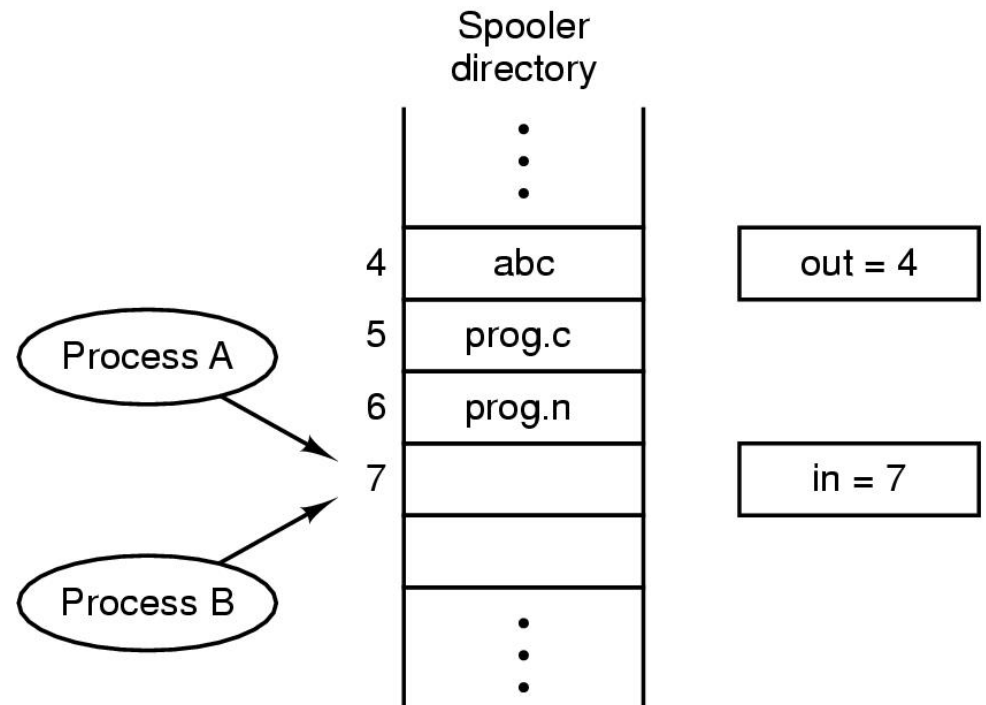
- **Race condition**
 - The result of the computation depends on the race between competing threads



Race Condition

- Two or more processes read/write some shared data, if the sequence of execution of the two processes makes a difference in the result of their concurrent execution, such situation is called **race conditions**

Without proper coordination, the result can be incorrect (and the result varies depending on which process runs when)!



Two processes want to access the spooler directory of a printer at the same time.

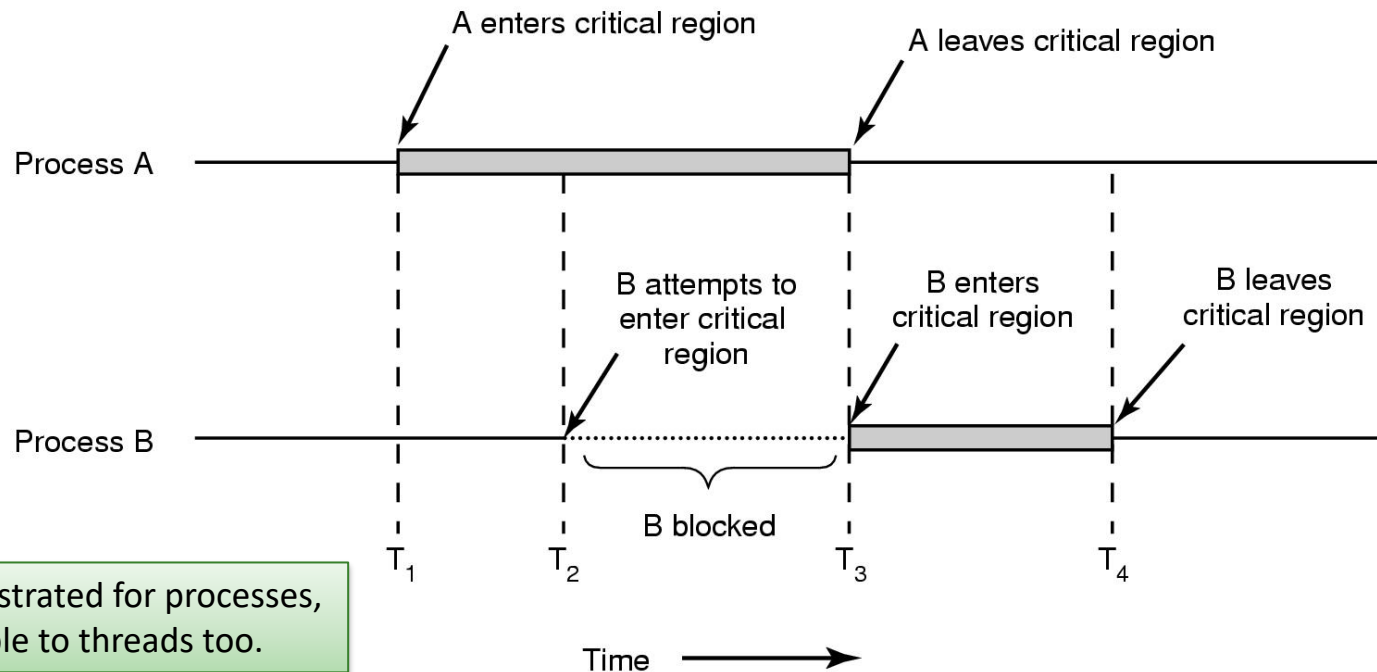


Thread-Safety(Cont.)

- Three levels of thread-safety:
https://en.wikipedia.org/wiki/Thread_safety
 - ❑ **Thread safe:** No race condition
 - ❑ **Conditionally safe:** shared resources are protected from race condition
 - ❑ **Not thread safe:** simultaneous accesses not supported
- Avoid Race Conditions (How to implement in Pthreads API):
 - ❑ Mutual exclusion
 - ❑ Thread-local storage: https://en.wikipedia.org/wiki/Thread-local_storage
 - ❑ Immutable Storage: https://en.wikipedia.org/wiki/Immutable_object
 - ❑ Atomic operation <https://en.wikipedia.org/wiki/Linearizability>
 - ❑ Re-entrancy [https://en.wikipedia.org/wiki/Reentrancy_\(computing\)](https://en.wikipedia.org/wiki/Reentrancy_(computing))

Solution to Race Conditions

- **Mutual exclusion** is needed to avoid race conditions and achieve correct results
 - Applies to any shared resources, e.g. memory, files, database, etc.



Mutual exclusion using critical regions.



Mutual Exclusion

- The code in the example corresponds to a **critical section**
 - A segment that must be executed by only one thread at any time
 - Test-and-update needs to be **atomic**
- Critical sections in Pthreads are implemented using **mutex locks (mutual exclusion locks)**
- Mutex-locks have two states: **locked** and **unlocked**
 - At any point of time, only one thread can lock a mutex lock
 - A lock is an atomic operation.
- A thread entering a critical section first tries to get a lock. It goes ahead when the lock is granted.



Mutual Exclusion (cont.)

- The Pthreads API provides the following functions for handling mutex-locks:

```
int pthread_mutex_lock (  
    pthread_mutex_t *mutex_lock);  
int pthread_mutex_unlock (  
    pthread_mutex_t *mutex_lock);  
int pthread_mutex_init (  
    pthread_mutex_t *mutex_lock,  
    const pthread_mutexattr_t *lock_attr);
```




Mutual Exclusion (cont.)

mutex_lock:

TSL REGISTER,MUTEX

CMP REGISTER,#0

JZE ok

CALL thread_yield

JMP mutex_lock

ok:

RET

| copy mutex to register and set mutex to 1
| was mutex zero?
| if it was zero, mutex was unlocked, so return
| mutex is busy; schedule another thread
| try again
| return to caller; critical region entered

Not busy waiting!

mutex_unlock:

MOVE MUTEX,#0

RET

| store a 0 in mutex
| return to caller

Possible implementation of mutex lock and mutex unlock.



Mutual Exclusion

- We can now write our previously incorrect code segment as:

```
pthread_mutex_t minimum_value_lock;  
...  
main() {  
    ....  
    pthread_mutex_init(&minimum_value_lock, NULL);  
    ....  
}  
void *find_min(void *list_ptr) {  
    ....  
    pthread_mutex_lock(&minimum_value_lock);  
    if (my_min < minimum_value)  
        minimum_value = my_min;  
    /* and unlock the mutex */  
    pthread_mutex_unlock(&minimum_value_lock);  
}
```

} Critical section



Outline

- Questions?
- Overview of programming models and thread basics
- The POSIX Threads
- Synchronization Primitives in Pthreads
 - Mutual Exclusion for Shared Variables
 - Producer-Consumer problem
 - Condition Variables for Synchronization
 - Controlling Thread and Synchronization Attributes
 - Composite Synchronization Constructs
- Compiling and running Pthreads program



Producer-Consumer Problem

- The producer-consumer relationship models a class of problems, with the following constraints
- The producer thread must not overwrite the shared buffer when the previous task has not been picked up by a consumer thread.
- The consumer threads must not pick up tasks until there is something present in the shared data structure.
- Producer-consumer relations are ubiquitous
- Needs synchronization between producer and consumer



The Producer-Consumer Problem

```
#define N 100
int count = 0;

void producer(void)
{
    int item;

    while (TRUE) {
        item = produce_item();
        if (count == N) sleep();
        insert_item(item);
        count = count + 1;
        if (count == 1) wakeup(consumer);
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {
        if (count == 0) sleep();
        item = remove_item();
        count = count - 1;
        if (count == N - 1) wakeup(producer);
        consume_item(item);
    }
}
```

/* number of slots in the buffer */
/* number of items in the buffer */

/* repeat forever */
/* generate next item */
/* if buffer is full, go to sleep */
/* put item in buffer */
/* increment count of items in buffer */
/* was buffer empty? */

/* repeat forever */
/* if buffer is empty, got to sleep */
/* take item out of buffer */
/* decrement count of items in buffer */
/* was buffer full? */
/* print item */

Code **NOT** protected by mutual exclusion! Can have **incorrect result!**

e.g. a context switch can occur here right after testing the count has a value of 0

The producer-consumer problem with a fatal race condition.



Producer-Consumer Using Locks

```
pthread_mutex_t task_queue_lock;
int task_available;

...
main() {
    ....
    task_available = 0;
    pthread_mutex_init(&task_queue_lock, NULL);
    ....
}

void *producer(void *producer_thread_data) {
    ....
    while (!done()) {
        inserted = 0;
        create_task(&my_task);
        while (inserted == 0) {
            pthread_mutex_lock(&task_queue_lock);
            if (task_available == 0) {
                insert_into_queue(my_task);
                task_available = 1;
                inserted = 1;
            }
            pthread_mutex_unlock(&task_queue_lock);
        }
    }
}
```

Assumes task queue
holds only one task

Signal the task availability

Release the lock



Producer-Consumer Using Locks

```
void *consumer(void *consumer_thread_data) {
    int extracted;
    struct task my_task;
    /* local data structure declarations */
    while (!done()) {
        extracted = 0;
        while (extracted == 0) {
            pthread_mutex_lock(&task_queue_lock);
            if (task_available == 1) {
                extract_from_queue(&my_task);
                task_available = 0;
                extracted = 1;
            }
            pthread_mutex_unlock(&task_queue_lock);
        }
        process_task(my_task);
    }
}
```

Signal the queue availability

Release the lock



Types of Mutexes

- Pthreads supports three types of mutexes - normal, recursive, and error-check.
- A **normal mutex** deadlocks if a thread that already has a lock tries a second lock on it.
- A **recursive mutex** allows a single thread to lock a mutex as many times as it wants. It simply increments a count on the number of locks. A lock is relinquished by a thread when the count becomes zero.
- An **error check mutex** reports an error when a thread with a lock tries to lock it again (as opposed to deadlocking in the first case, or granting the lock, as in the second case).
- The type of the mutex can be set in the **attributes object** before it is passed at time of initialization.



Overheads of Locking

- Locks represent **serialization points** since critical sections must be executed by threads one after the other.
- Encapsulating **large segments of the program within locks** can lead to significant performance degradation.
- It is often possible to reduce the idling overhead associated with locks using an alternate nonblocking function, **pthread_mutex_trylock**.

```
int pthread_mutex_trylock (  
    pthread_mutex_t *mutex_lock);
```
- Returns EBUSY when failed, allow the thread to do other work and poll the mutex for a lock later
- Example in the textbook



Avoiding deadlocks

- Deadlocks can occur when using more than one mutex and to lock multiple mutex from multiple threads

Thread 0

```
pthread_mutex_lock(&mutex_a);  
pthread_mutex_lock(&mutex_b);
```

Thread 1

```
pthread_mutex_lock(&mutex_b);  
pthread_mutex_lock(&mutex_a);
```

- Avoiding deadlock with **hierarchical locking**
 - Order locks
 - All threads follow the same order to lock multiple locks
- Avoiding deadlock with **nonblocking locking**
 - After locking the first mutex, use `pthread_mutex_trylock` to lock additional mutexes
 - If an attempt fails, release all mutexes and start again



Conditions for Deadlocks

- **Mutual exclusion condition**: each lock (or resource in general) is either currently assigned to exactly one process/thread or is available
- **Hold and wait condition**: processes/threads currently holding resources that were granted earlier can request new resources
- **No preemption condition**: resources previously granted cannot be forcibly taken away from a process/thread (must be voluntarily released)
- **Circular wait condition**: there is a circular list of two or more processes/threads, each of which is waiting for a resource held by another process/thread

All four conditions must be present for a deadlock to occur. If one of them is absent, no deadlock is possible. Hierarchical locking breaks the “Circular wait condition”.



Outline

- Questions?
- Overview of programming models and thread basics
- The POSIX Threads
- Synchronization Primitives in Pthreads
 - Mutual Exclusion for Shared Variables
 - Producer-Consumer problem
 - Condition Variables for Synchronization
 - Controlling Thread and Synchronization Attributes
 - Composite Synchronization Constructs
- Compiling and running Pthreads program



Condition Variables for Synchronization

- Idling overhead for `pthread_mutex_lock()`
- Polling overhead for `pthread_mutex_trylock()`
- A **condition variable** allows a thread to suspend its execution until specified data reaches a predefined state
 - ❑ Interrupt driven mechanism instead of a polled mechanism



Condition Variables for Synchronization

- A condition variable is associated with a **predicate**. When the predicate becomes true, the condition variable is used to signal one or more threads waiting on the condition.
- A condition variable **always has a mutex associated with it**. A thread locks this mutex and tests the predicate defined on the shared variable.
- A single condition variable may be associated with more than one predicate (**difficult to debug, discouraged**)



Condition Variables for Synchronization

- Pthreads provides the following functions for condition variables:

```
int pthread_cond_wait(pthread_cond_t *cond,  
    pthread_mutex_t *mutex);  
int pthread_cond_timedwait(pthread_cond_t *cond,  
    pthread_mutex_t *mutex,  
    const struct timespec *abstime);  
int pthread_cond_signal(pthread_cond_t *cond);  
int pthread_cond_broadcast(pthread_cond_t *cond);  
int pthread_cond_init(pthread_cond_t *cond,  
    const pthread_condattr_t *attr);  
int pthread_cond_destroy(pthread_cond_t *cond);
```



Producer-Consumer Using Condition Variables

```
pthread_cond_t cond_queue_empty, cond_queue_full;
pthread_mutex_t task_queue_cond_lock;
int task_available;
/* other data structures here */
main() {
    /* declarations and initializations */
    task_available = 0;
    pthread_init();
    pthread_cond_init(&cond_queue_empty, NULL);
    pthread_cond_init(&cond_queue_full, NULL);
    pthread_mutex_init(&task_queue_cond_lock, NULL);
    /* create and join producer and consumer threads */
}
```




Producer-Consumer Using Condition Variables

```
void *producer(void *producer_thread_data) {
    int inserted;
    while (!done()) {
        create_task();
        pthread_mutex_lock(&task_queue_cond_lock);
        while (task_available == 1)
            pthread_cond_wait(&cond_queue_empty,
                             &task_queue_cond_lock);
        insert_into_queue();
        task_available = 1;
        pthread_cond_signal(&cond_queue_full);
        pthread_mutex_unlock(&task_queue_cond_lock);
    }
}
```



Producer-Consumer Using Condition Variables

```
void *consumer(void *consumer_thread_data) {  
    while (!done()) {  
        pthread_mutex_lock(&task_queue_cond_lock);  
        while (task_available == 0)  
            pthread_cond_wait(&cond_queue_full,  
                             &task_queue_cond_lock);  
        my_task = extract_from_queue();  
        task_available = 0;  
        pthread_cond_signal(&cond_queue_empty);  
        pthread_mutex_unlock(&task_queue_cond_lock);  
        process_task(my_task);  
    }  
}
```



Readings

- Reference book ITPC – Chapter 7
- POSIX Threads Programming, by Blaise Barney, Lawrence Livermore National Laboratory: <https://computing.llnl.gov/tutorials/pthreads/>



Questions?

Questions/Suggestions/Comments are always welcome!

Write me: yong.chen@ttu.edu

Call me: 806-834-0284

See me: ENGCTR 315

If you write me an email for this class, please start the email subject with [CS4379] or [CS5379].