



CS4379: Parallel and Concurrent Programming

CS5379: Parallel Processing

Lecture 16

Guest Lecture by Wei Zhang

X-Spirit.Zhang@ttu.edu

Dr. Yong Chen

Associate Professor

Computer Science Department

Texas Tech University



Course Info

- **Lecture Time:** TR, 12:30-1:50
- **Lecture Location:** ECE 217
- **Sessions:** CS4379-001, CS4379-002, CS5379-001, CS5379-D01
- **Instructor:** Yong Chen, Ph.D., Associate Professor
- **Email:** yong.chen@ttu.edu
- **Phone:** 806-834-0284
- **Office:** Engineering Center 315
- **Office Hours:** 2-4 p.m. on Wed., or by appointment
- **TA:** Mr. Ghazanfar Ali, Ghazanfar.Ali@ttu.edu
- **TA Office hours:** Tue. and Fri., 2-3 p.m., or by appointment
- **TA Office:** EC 201 A
- **More info:**
 - <http://www.myweb.ttu.edu/yonchen>
 - <http://discl.cs.ttu.edu>; <http://cac.ttu.edu/>; <http://nsfcac.org>



Outline

- Questions?
- Overview of programming models and thread basics
- The POSIX Threads
- Synchronization Primitives in Pthreads
 - Mutual Exclusion for Shared Variables
 - Producer-Consumer problem
 - Condition Variables for Synchronization
 - Controlling Thread and Synchronization Attributes
 - Composite Synchronization Constructs
- Compiling and running Pthreads program



Controlling Thread and Synchronization

Attributes

- The Pthreads API allows a programmer to change the default attributes of entities using **attributes objects**.
- An attributes object is a data-structure that describes entity (thread, mutex, condition variable) properties.
- Once these properties are set, the attributes object can be passed to the method initializing the entity.
- Separates semantics and implementation
- Enhances modularity, readability, and ease of modification.



Attribute Objects for Threads

- Use `pthread_attr_init` to create an attribute object.
- Individual properties associated with the attributes object can be changed using the following functions:

`pthread_attr_setdetachstate,`
`pthread_attr_setstacksize,`
`pthread_attr_setinheritsched,`
`pthread_attr_setschedpolicy,` and
`pthread_attr_setschedparam`

- Provide fine control



Attribute Objects for Mutexes

- Initialize the attributes object using function:

`pthread_mutexattr_init.`

- The function `pthread_mutexattr_settype` can be used for setting the type of mutex specified by the mutex attributes object.

```
pthread_mutexattr_settype (  
pthread_mutexattr_t *attr,  
int type);
```

- Here, `type` specifies the type of the mutex and can take one of:
 - ❑ `PTHREAD_MUTEX_NORMAL_NP`
 - ❑ `PTHREAD_MUTEX_RECURSIVE_NP`
 - ❑ `PTHREAD_MUTEX_ERRORCHECK_NP`



Composite Synchronization Constructs

- By design, Pthreads provide support for a basic set of operations.
- Higher level constructs can be built using basic synchronization constructs
- We discuss two such constructs - **read-write locks** and **barriers**.



Read-Write Locks

- In many applications, a data structure is **read frequently** but written **infrequently**. For such applications, we should use **read-write locks**.
- Example: <http://ianfinlayson.net/class/cpsc425/notes/08-read-write>
- Initially unlocked, read/write lock can be obtained
- If read locked
 - Read locks granted, reader counter incremented
 - Write lock performs a condition wait, pending writer counter incremented
- If write locked
 - Read/write locks perform condition wait, pending writer counter incremented
- Reading: https://en.wikipedia.org/wiki/Readers–writer_lock



Read-Write Locks

- The lock data type `mylib_rwlock_t` holds the following:
 - ❑ A count of the number of readers, writers
 - ❑ A condition variable `readers_proceed` that is signaled when readers can proceed,
 - ❑ A condition variable `writer_proceed` that is signaled when one of the writers can proceed,
 - ❑ A count `pending_writers` of pending writers, and
 - ❑ A mutex `read_write_lock` associated with the shared data structure



Read-Write Locks

```
void mylib_rwlock_rlock(mylib_rwlock_t *l) {
    if there is a write lock or pending writers, perform
    condition wait.. else increment count of readers and grant
    read lock
}

void mylib_rwlock_wlock(mylib_rwlock_t *l) {
    if there are readers or writers, increment pending writers
    count and wait. On being woken, decrement pending writers
    count and increment writer count
}

void mylib_rwlock_unlock(mylib_rwlock_t *l) {
    if there is a write lock then unlock, else if there are read
    locks, decrement count of read locks. If the count is 0 and
    there is a pending writer, let it through, else if there are
    pending readers, let them all go through
}
```




Barriers (1)


- One commonly used synchronization mechanism in IPC
- A process/thread is blocked at barrier till **all processes/threads** participating in the barrier reach it
- Useful for **applications with phases** and have the rule that no process/thread may proceed into the next phase until all processes/threads are ready to proceed
 - E.g. iterations of calculations with many time steps in scientific apps




Barriers (2)



(a)



(b)



(c)

- Use of a barrier. (a) Processes approaching a barrier.
(b) All processes but one blocked at the barrier.
(c) When the last process arrives at the barrier, all of them are let through.



Barriers Implementation: a Linear Barrier

- Barriers can be implemented using a counter, a mutex and a condition variable.
- A single integer is used to keep track of the number of threads that have reached the barrier.
- If the count is less than the total number of threads, the threads execute a condition wait.
- The last thread entering (and setting the count to the number of threads) wakes up all the threads using a condition broadcast.




Barriers Implementation: a **Logarithmic Barrier**

- In the prior implementation, threads enter the barrier and stay until the broadcast signal releases them, and threads released one by one
- A logarithmic/tree barrier implementation pairs threads and uses $n/2$ condition variables, more efficient than a linear barrier
- See the reference book (ITPC) for detailed codes



Log Barrier v.s. Linear Barrier



- Execution time of 1000 sequential and logarithmic barriers as a function of number of threads on a 32 processor SGI Origin 2000.



Outline

- Questions?
- Overview of programming models and thread basics
- The POSIX Threads
- Synchronization Primitives in Pthreads
- Compiling and running Pthreads program



Sample Source Codes

- Once you log in HPCC cluster, you can check out source codes with executing the following command:

```
git clone https://discl.cs.ttu.edu/gitlab/yongchen/cs4379cs5379.git
```



Compiling Pthreads programs

- Load GNU Compiler module

- ❑ `$ module load gnu7/7.3.0`

- Compiling

- ❑ `$ gcc -lpthread (or g++ -lpthread)`

- ❑ E.g. `$ gcc -Wall -o gauss_pthreads -O2 -lpthread gauss_pthreads.c`

- ❑ Or, `$ gcc -o hello-pthread -lpthread hello-pthread.c`

- Automate compiling using Makefile

- ❑ Edit a Makefile

- ❑ `$ make`



Makefile – Automating Compilation

constant

■ Makefile:

- ❑ A simple way to organize code compilation
- ❑ `$@` says to put the output of the compilation in the file named on the left side of the `:`
- ❑ `^` means right side of `:`
- ❑ `clean:` tells make to clean the object files

```
CC=gcc
OBJ=gauss_threads.o

%.o: %.c
    $(CC) -c -o $@ $^

gauss_threads: $(OBJ)
    $(CC) -o $@ $^

clean:
    rm -f *.o
```

tab



Makefile – Automating Compilation

- For how to write a Makefile, such as write rules and targets, please see:
 - ❑ <http://www.gnu.org/software/make/manual/make.pdf> or
 - ❑ http://www.gnu.org/software/make/manual/html_node/index.html
 - ❑ <https://makefiletutorial.com>
 - ❑ <https://www.tutorialspoint.com/makefile/index.htm>
- More info: <http://www.gnu.org/software/make/>



Debugging Pthread Program

- Prerequisite: set `-g` option for gcc
- GDB/CGDB: (`-ggdb`)

```
set print thread-events    # prints out thread start and exit events
info threads               # list all existing threads in program
                           # the gdb threadno is the first value listed
                           # the thread that hit the break point is *'ed
thread threadno            # switch to thread threadno's context
                           # (see its stack when type where, for example)
break [where] thread [threadno] # set a breakpoint at [where] just for
                                # thread threadno

# apply the gdb command to all or a subset of threads
thread apply [threadno|all] command
```



Running Pthreads programs

- Create job submission script, e.g.

```
#!/bin/sh
#$ -V
#$ -cwd
#$ -S /bin/bash
#$ -N Pthreads_Test_Job
#$ -o $JOB_NAME.o$JOB_ID
#$ -e $JOB_NAME.e$JOB_ID
#$ -q omni
#$ -pe sm 36
#$ -l h_vmem=5.3G
#$ -l h_rt=48:00:00
#$ -P quanah
```

```
./hello-pthreads
```

- Submit job
 - ❑ **qsub** <job submission script>
- Check job status
 - ❑ Command: **qstat**



Checking Output and Debugging Failed Jobs


- Job output
 - ❑ Standard: \$JOB_NAME.o\$JOB_ID
 - ❑ Error: \$JOB_NAME.e\$JOB_ID

- When debugging:
 - ❑ Check the output files for errors
 - ❑ Check the output of qacct -j <job_ID>
 - failed
 - exit_status
 - maxvmem
 - start_time & end_time (<runtime limit)
 - low



More about Pthreads

- IEEE Std 1003.1, 2004 Edition
- <http://pubs.opengroup.org/onlinepubs/009695399/basedefs/pthread.h.html>
- “Programming with POSIX Threads”, by David R. Butenhof, ISBN-10: 0201633922, ISBN-13: 978-0201633924





Readings

- Reference book ITPC – Chapter 7
- POSIX Threads Programming, by Blaise Barney, Lawrence Livermore National Laboratory: <https://computing.llnl.gov/tutorials/pthreads/>



Questions?

Questions/Suggestions/Comments are always welcome!

Write me: yong.chen@ttu.edu

Call me: 806-834-0284

See me: ENGCTR 315

If you write me an email for this class, please start the email subject with [CS4379] or [CS5379].