



CS4379: Parallel and Concurrent Programming

CS5379: Parallel Processing

Lecture 17

Dr. Yong Chen

Associate Professor

Computer Science Department

Texas Tech University



Lecture Video

- Please view the lecture video either from Teams or from the below link:
- <https://texastechuniversity.sharepoint.com/sites/CS4379-CS5379/Shared%20Documents/General/Lecture17.mp4>



Course Info

- **Lecture Time:** TR, 12:30-1:50
- **Lecture Location:** ECE 217
- **Sessions:** CS4379-001, CS4379-002, CS5379-001, CS5379-D01
- **Instructor:** Yong Chen, Ph.D., Associate Professor
- **Email:** yong.chen@ttu.edu
- **Phone:** 806-834-0284
- **Office:** Engineering Center 315
- **Office Hours:** 2-4 p.m. on Wed., or by appointment
- **TA:** Mr. Ghazanfar Ali, Ghazanfar.Ali@ttu.edu
- **TA Office hours:** Tue. and Fri., 2-3 p.m., or by appointment
- **TA Office:** EC 201 A
- **More info:**
 - <http://www.myweb.ttu.edu/yonchen>
 - <http://discl.cs.ttu.edu>; <http://cac.ttu.edu/>; <http://nsfcac.org>



Announcements/Reminders

- Recorded lectures for the rest of semester – due to mixed sessions and distance section requires viewing lectures at a flexible time
- Remaining quizzes will be scheduled at class time and will be announced in advance
- Skype meeting link for my office hours (Wed. 2 – 4 p.m.)
 - <https://meet.ttu.edu/yong.chen/60DC09T2>
 - Join by Phone
 - +1 (806) 834-4888,, 64158017#
 - +1 (855) 834-4888 - Toll-free,, 64158017#
 - Conference ID: 64158017
- Please be safe and stay well



Outline

- Questions?
- OpenMP programming model
- Specifying concurrent tasks in OpenMP




OpenMP: a Standard for Directive Based Parallel Programming

- Pthreads still need explicit management of threads, low-level
- High-level constructs/directives desired
- OpenMP provides a **directive-based API** that can be used with FORTRAN, C, and C++ for programming shared address space machines.
 - OpenMP: Open Multi-Processing
- OpenMP directives provide **support for concurrency, synchronization, and data handling** while avoiding the need for explicitly setting up threads, distributing tasks, managing mutex locks, etc.



OpenMP programming model

- OpenMP uses the **fork-join model** of parallel execution:
 - **Begin as a single thread**: the master thread and executes sequentially until the first parallel region construct is encountered
 - **FORK**: the master thread then creates a team of parallel threads
 - The statements in the program that are enclosed by the parallel region construct are then executed in parallel among the various team threads.
 - **JOIN**: When the team threads complete, they synchronize and terminate, leaving only the master thread.





OpenMP Programming Model

- OpenMP directives in C and C++ are based on the **#pragma compiler directives**
- A directive consists of a directive name followed by clauses.
`#pragma omp directive [clause list]`
- OpenMP programs execute serially until they encounter the `parallel` directive, which creates a group of threads.
`#pragma omp parallel [clause list]`
`/* structured block */`
- The main thread that encounters the `parallel` directive becomes the *master* of this group of threads and is assigned the thread id 0 within the group.



OpenMP Programming Model

- The clause list is used to specify conditional parallelization, number of threads, and data handling.
 - **Conditional Parallelization:** The clause `if (scalar expression)` determines whether the parallel construct results in creation of threads.
 - **Degree of Concurrency:** The clause `num_threads(integer expression)` specifies the number of threads that are created.

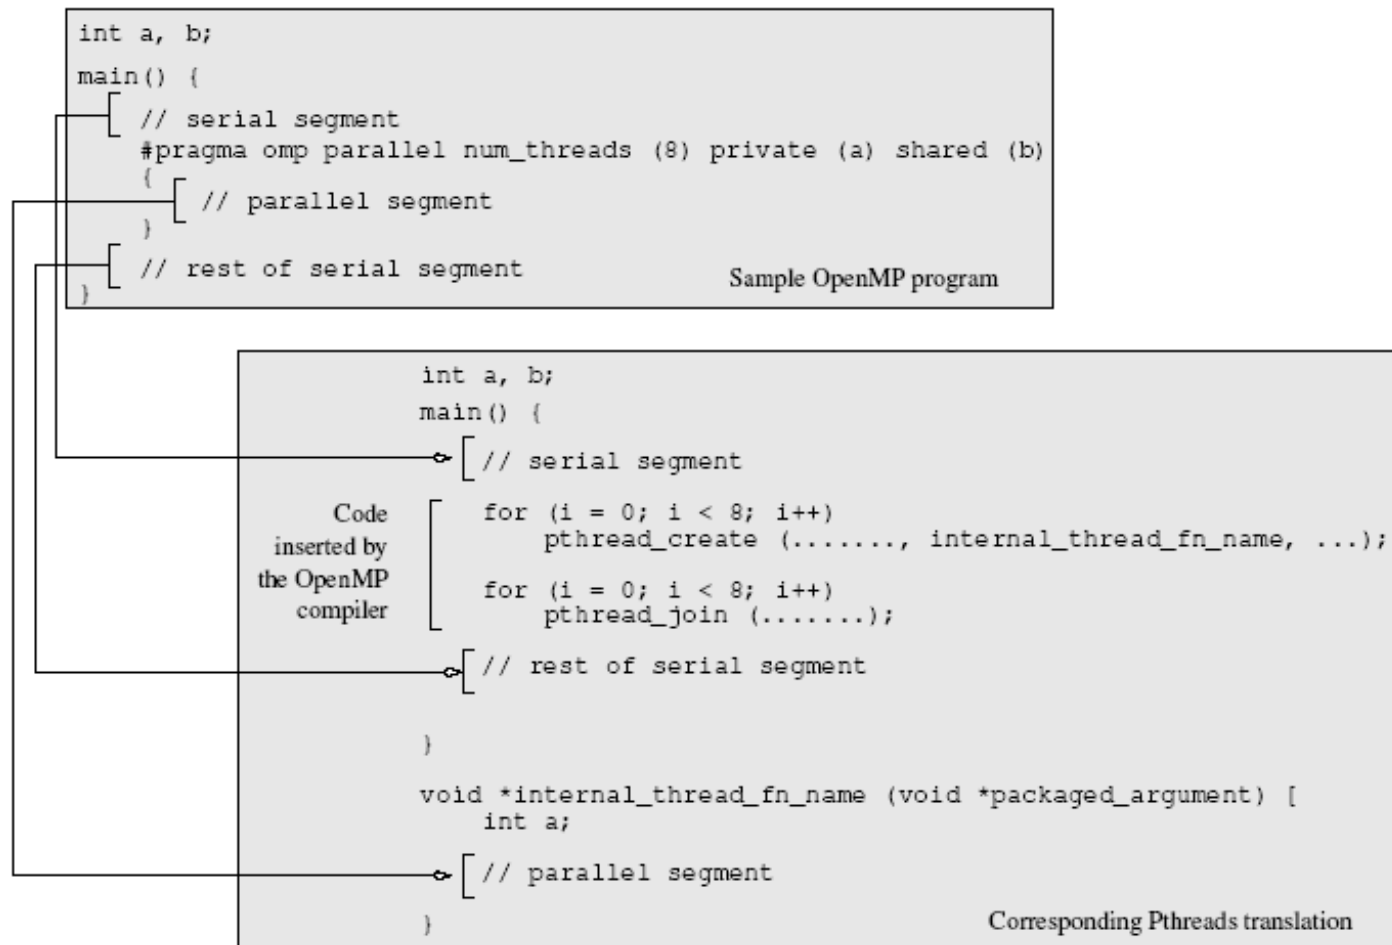


OpenMP Programming Model

- The clause list is used to specify conditional parallelization, number of threads, and data handling (cont.)
 - ❑ **Data Handling:** The clause `private` (variable list) indicates variables local to each thread.
 - ❑ The clause `firstprivate` (variable list) is similar to the `private`, except values of variables are initialized to corresponding values before the parallel directive.
 - ❑ The clause `shared` (variable list) indicates that variables are shared across all the threads.
 - ❑ The default state of a variable is specified by the clause `default`, e.g. `default(shared)`, or `default(private)`



OpenMP Programming Model



- A sample OpenMP program along with its Pthreads translation that might be performed by an OpenMP compiler.



OpenMP Programming Model

```
#pragma omp parallel if (is_parallel == 1) num_threads(8) \  
    private (a) shared (b) firstprivate(c) {  
    /* structured block */  
}
```

- If the value of the variable `is_parallel` equals one, eight threads are created.
- Each of these threads gets private copies of variables `a` and `c`, and shares a single value of variable `b`.
- The value of each copy of `c` is initialized to the value of `c` before the parallel directive.



Reduction Clause in OpenMP

- The `reduction` clause specifies how multiple local copies of a variable at different threads are combined into a single copy at the master when threads exit.
- The usage of the `reduction` clause is `reduction (operator: variable list)`.
- The variables in the list are implicitly specified as being private to threads.
- The operator can be one of `+`, `*`, `-`, `&`, `|`, `^`, `&&`, and `||`.

```
#pragma omp parallel reduction(+: sum) num_threads(8) {  
/* compute local sums here */  
}  
/*sum here contains sum of all local instances of sums */
```



OpenMP Programming: Example

```
/* *****  
An OpenMP version of a threaded program to compute Pi.  
***** */  
#pragma omp parallel default(private) shared (npoints) \  
    reduction(+: sum) num_threads(8)  
{  
    num_threads = omp_get_num_threads();  
    sample_points_per_thread = npoints / num_threads;  
    sum = 0;  
    for (i = 0; i < sample_points_per_thread; i++) {  
        rand_no_x = (double)(rand_r(&seed)) / (double)((2<<14) -  
            1);  
        rand_no_y = (double)(rand_r(&seed)) / (double)((2<<14) -  
            1);  
        if ((rand_no_x * rand_no_x +  
            rand_no_y * rand_no_y) < 1)  
            sum++;  
    }  
}
```

← OpenMP library functions



Outline

- Questions?
- OpenMP programming model
- Specifying concurrent tasks in OpenMP



Specifying Concurrent Tasks in OpenMP

- The `parallel` directive can be used in conjunction with other directives to **specify concurrency across iterations and tasks**.
- OpenMP provides **two directives: `for` and `sections`** - to specify concurrent iterations and tasks.



Specifying Concurrent Tasks in OpenMP

- The **for directive** is used to split parallel iteration spaces across threads. The general form of a for directive is as follows:

```
#pragma omp for [clause list]
/* for loop */
```

- The clauses that can be used in this context include: **private**, **firstprivate**, **reduction**, **schedule**, **nowait**.



Specifying Concurrent Tasks in OpenMP: for Directive Example (compute Pi)

```
#pragma omp parallel default(private) shared (npoints) \  
    reduction(+: sum) num_threads(8)  
{  
    sum = 0;  
    #pragma omp for  
    for (i = 0; i < npoints; i++) {  
        rand_no_x =(double)(rand_r(&seed))/(double)((2<<14)-1);  
        rand_no_y =(double)(rand_r(&seed))/(double)((2<<14)-1);  
        if ((rand_no_x * rand_no_x +  
            rand_no_y * rand_no_y) < 1)  
            sum++;  
    }  
}
```

Simple OpenMP programming:
only two directives



Assigning Iterations to Threads

- The `schedule` clause of the `for` directive deals with the assignment of iterations to threads.
- The general form of the `schedule` directive is `schedule(scheduling_class[, parameter])`.
- OpenMP supports four scheduling classes: `static`, `dynamic`, `guided`, and `runtime`.




Assigning Iterations to Threads: Example

```
/* static scheduling of matrix multiplication loops */
#pragma omp parallel default(private) shared (a, b, c, dim) \
    num_threads(4)
#pragma omp for schedule(static)
for (i = 0; i < dim; i++) {
    for (j = 0; j < dim; j++) {
        c(i,j) = 0;
        for (k = 0; k < dim; k++) {
            c(i,j) += a(i, k) * b(k, j);
        }
    }
}
```

`schedule(static[, chunk-size])` splits the iteration space into equal chunks of size `chunk-size` and assigns them to threads in a round-robin fashion

Assigning Iterations to Threads: Example



`schedule(static)` `schedule(static, 16)` Each for loop is parallelized and `schedule(static)`

- Three different schedules using the static scheduling class of OpenMP.



Scheduling Strategies

- **Dynamic**: assigned to threads when they become idle
 - ❑ `schedule(dynamic[, chunk-size])`.
 - ❑ Default to single iteration per chunk
- **Guided**: chunk size is dynamically changed toward the completion of the computation
 - ❑ To avoid load imbalancing and idling
- **Runtime**: delay scheduling decision until runtime, not in the code
 - ❑ Environment variable `OMP_SCHEDULE` determines the scheduling class and chunk size



Synchronization Across Multiple for Directives

- It is often desirable to have a sequence of `for`-directives within a parallel construct that do not execute an implicit barrier at the end of each `for` directive.
- OpenMP provides a clause - `nowait`, which can be used with a `for` directive.



Synchronization Across Multiple for Directives

```
#pragma omp parallel
{
    #pragma omp for nowait
        for (i = 0; i < nmax; i++)
            if (isEqual(name, current_list[i])
                processCurrentName(name);
    #pragma omp for
        for (i = 0; i < mmax; i++)
            if (isEqual(name, past_list[i])
                processPastName(name);
}
```




The sections Directive

- OpenMP supports **non-iterative parallel task assignment** using the `sections` directive.
- The general form of the `sections` directive is as follows:

```
#pragma omp sections [clause list]
{
    [#pragma omp section
        /* structured block */
    ]
    [#pragma omp section
        /* structured block */
    ]
    ...
}
```



The sections Directive: Example

```
#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
        {
            taskA();
        }
        #pragma omp section
        {
            taskB();
        }
        #pragma omp section
        {
            taskC();
        }
    }
}
```



Merging Directives

- Without parallel directive specified, the `for` and `sections` directives would execute serially
- Consequently, `for` and `sections` directives are generally preceded by the parallel directive
- OpenMP allows the programmer to merge the parallel directives to `parallel for` and `parallel sections`, respectively
- The clause list for the merged directive can be from the clause lists of either the parallel or `for / sections` directives.



Merging Directives Example


```
1  #pragma omp parallel default (private) shared (n)
2  {
3      #pragma omp for
4      for (i = 0 < i < n; i++) {
5          /* body of parallel for loop */
6      }
7  }
```

is identical to:


```
1  #pragma omp parallel for default (private) shared (n)
2  {
3      for (i = 0 < i < n; i++) {
4          /* body of parallel for loop */
5      }
6  }
```



Merging Directives



parallel for



parallel sections



Readings

- Reference book ITPC – Chapter 7, 7.10
- OpenMP Programming, by Blaise Barney, Lawrence Livermore National Laboratory: <https://computing.llnl.gov/tutorials/openMP/>
- OpenMP 5.0 Complete Specifications
 - <https://www.openmp.org/wp-content/uploads/OpenMP-API-Specification-5.0.pdf>



Questions?

Questions/Suggestions/Comments are always welcome!

Write me: yong.chen@ttu.edu

Call me: 806-834-0284

See me: ENGCTR 315

If you write me an email for this class, please start the email subject with [CS4379] or [CS5379].