



CS4379: Parallel and Concurrent Programming

CS5379: Parallel Processing

Lecture 18

Dr. Yong Chen

Associate Professor

Computer Science Department

Texas Tech University



Lecture Video

- Please view the lecture video either from Teams or from the below link:
- <https://texastechuniversity.sharepoint.com/sites/CS4379-CS5379/Shared%20Documents/General/Lecture18.mp4>



Course Info

- **Lecture Time:** TR, 12:30-1:50
- **Lecture Location:** ECE 217
- **Sessions:** CS4379-001, CS4379-002, CS5379-001, CS5379-D01
- **Instructor:** Yong Chen, Ph.D., Associate Professor
- **Email:** yong.chen@ttu.edu
- **Phone:** 806-834-0284
- **Office:** Engineering Center 315
- **Office Hours:** 2-4 p.m. on Wed., or by appointment
- **TA:** Mr. Ghazanfar Ali, Ghazanfar.Ali@ttu.edu
- **TA Office hours:** Tue. and Fri., 2-3 p.m., or by appointment
- **TA Office:** EC 201 A
- **More info:**
 - <http://www.myweb.ttu.edu/yonchen>
 - <http://discl.cs.ttu.edu>; <http://cac.ttu.edu/>; <http://nsfcac.org>



Outline

- Questions?
- Synchronization constructs in OpenMP
- OpenMP runtime library functions
- Environment variables
- Compiling, running, and demos



Synchronization Constructs in OpenMP

- OpenMP provides a variety of synchronization constructs:

```
#pragma omp barrier
```

```
#pragma omp single [clause list]  
    structured block
```

```
#pragma omp master  
    structured block
```

```
#pragma omp critical [(name)]  
    structured block
```

```
#pragma omp atomic  
    expression
```



Synchronization Point: The barrier Directive

- One of the most frequently used synchronization primitives

```
#pragma omp barrier
```

- On encountering this directive, all threads in a team wait until others have caught up, and then release
- When used with nested parallel directives, the barrier directive binds to the closest parallel directive



Single Thread Executions: The single and master Directives


- A single directive specifies a structured block that is executed by a single (arbitrary) thread.

```
#pragma omp single [clause list]  
    structured block
```

- E.g. one thread compute the mean by dividing this global sum by the number of entries in the list

- The master directive is a specialization of the single directive in which only the master thread executes the structured block

```
#pragma omp master  
    structured block
```





Critical Sections: The critical Directive

- OpenMP provides a critical directive for implementing critical regions

```
#pragma omp critical [(name)]  
    structured block
```

- The use of name can identify a critical region and allows different threads to execute different critical regions
- OpenMP can also have explicit lock management



Critical Sections: The atomic Directive

- Often, a critical section consists simply of an update to a **single memory location**, e.g. incrementing or adding to an integer
- OpenMP provides another directive, **atomic**, for such atomic update to a memory location.
- Applied only to the single statement that immediately follows it

```
#pragma omp atomic
    count++;
```

- All atomic directives can be replaced by critical directives
- However, the availability of atomic hardware instructions may optimize the performance of the program



OpenMP Runtime Library Functions

- In addition to directives, the OpenMP standard defines an API for library calls that control the execution of threaded programs
 - ❑ Query the number of threads/processors, set number of threads to use
 - ❑ General purpose locking routines
 - ❑ Portable wall clock timing routines
 - ❑ Set execution environment functions: nested parallelism, dynamic adjustment of threads.

- It is necessary to specify the `include file omp.h`

The OpenMP specification is your best friend and reference.



Controlling Number of Threads and Processors

- Functions relate to the concurrency
- `void omp_set_num_threads(int num_threads);`
 - Sets the number of threads that will be used in the next parallel region (a positive integer)
 - This routine can only be called from the serial portions of the code
 - This call has precedence over the OMP_NUM_THREADS environment variable
- `int omp_get_num_threads();`
 - Returns the number of threads that are currently in the team executing the parallel region from which it is called.
 - If this call is made from a serial portion of the program, it will return 1



Example

- \$ gcc -fopenmp

- omp_setnumthreads.c

- \$./a.out

```
#include <stdio.h>
#include <omp.h>

int main()
{
    printf("%d\n", omp_get_num_threads( ));
    omp_set_num_threads(4);
    printf("%d\n", omp_get_num_threads( ));
    #pragma omp parallel
        #pragma omp master
        {
            printf("%d\n", omp_get_num_threads( ));
        }

    printf("%d\n", omp_get_num_threads( ));

    #pragma omp parallel num_threads(3)
        #pragma omp master
        {
            printf("%d\n", omp_get_num_threads( ));
        }

    printf("%d\n", omp_get_num_threads( ));
}
```



Controlling Number of Threads and Processors

- `int omp_get_max_threads();`
 - Returns the maximum number of threads that could possibly be created
 - May be called from both serial and parallel regions of code
- `int omp_get_thread_num();`
 - Returns the thread number of the thread, within the team
 - This number will be between 0 and OMP_GET_NUM_THREADS-1. The master thread of the team is thread 0
 - If called from a nested parallel region, or a serial region, return 0
- `int omp_get_num_procs();`
 - Returns the number of processors that are available to the program.
- `int omp_in_parallel();`
 - Determine if the section of code which is executing is parallel (non-zero) or not (zero)



Controlling and monitoring thread creation

- `void omp_set_dynamic(int dynamic_threads);`
 - Enables or disables dynamic adjustment of the number of threads available for execution of parallel regions
 - If the value evaluates to zero, disabled, otherwise it is enabled
 - Has precedence over the OMP_DYNAMIC environment variable.
 - Must be called from a serial section of the program
- `int omp_get_dynamic();`
 - Query if dynamic thread adjustment is enabled (non-zero) or not
- `void omp_set_nested(int nested);`
 - Enable (non-zero) or disable (zero) nested parallelism
 - Has precedence over the OMP_NESTED environment variable
- `int omp_get_nested();`
 - Determine if nested parallelism is enabled or not.



Mutual Exclusion

- `void omp_init_lock(omp_lock_t *lock);`
 - Initialize lock before using it
 - The lock data structure in OpenMP is of type `omp_lock_t`
- `void omp_destroy_lock(omp_lock_t *lock);`
- `void omp_set_lock(omp_lock_t *lock);`
 - Forces the executing thread to wait until the specified lock is available
 - A thread is granted ownership of a lock when it becomes available
- `void omp_unset_lock(omp_lock_t *lock);`
- `int omp_test_lock(omp_lock_t *lock);`
 - Attempt to set a lock but does not block if the lock is unavailable
 - If the function returns a non-zero value, the lock has been successfully set, otherwise the lock is currently owned by another thread



Timing Routines

- Get wall-clock time: a double precision value equal to the number of seconds since the initial value of the OS real-time clock.
 - ❑ `double omp_get_wtime(void)`
 - ❑ Provides a portable wall clock timing routine
 - ❑ Usually used in "pairs" with the value of the first call subtracted from the value of the second call to obtain the elapsed seconds

- Get wall-clock time precision:
 - ❑ `double omp_get_wtick(void)`
 - ❑ Get timer precision
 - ❑ Returns a double-precision floating point value equal to the number of seconds between successive clock ticks



Environment Variables in OpenMP

- **OMP_NUM_THREADS:**
 - ❑ This environment variable specifies the default number of threads created upon entering a parallel region
 - ❑ `csh: setenv OMP_NUM_THREADS 8`
 - ❑ `Bash: export OMP_NUM_THREADS=8`
- **OMP_DYNAMIC:**
 - ❑ Enables or disables dynamic adjustment of the number of threads
 - ❑ `csh: setenv OMP_DYNAMIC TRUE`
- **OMP_NESTED**
 - ❑ Enables or disables nested parallelism
- **OMP_SCHEDULE:**
 - ❑ Applies only to `for`, `parallel for` directives which schedule set to `RUNTIME`
 - ❑ `setenv OMP_SCHEDULE "dynamic"`
- All environment variable names are uppercase. The values assigned to them are not case sensitive



Explicit Threads versus Directive Based Programming

- Directives layered on top of threads facilitate a variety of thread-related tasks.
- Ease the tasks of creating/joining threads, setting up arguments to threads, partitioning iteration spaces, etc.
- There are **some drawbacks** to using directives as well.
 - Explicit threading makes data exchange more apparent, alleviating some of the overheads from data movement.
 - Explicit threading also provides a richer API in the form of condition waits, locks of different types, and increased flexibility for building composite synchronization operations.
 - Since explicit threading is used more widely than OpenMP, tools and support for Pthreads programs are easier to find.



Outline

- Questions?
- Synchronization Constructs in OpenMP
- OpenMP runtime library functions
- Environment variables
- Compiling, running, and demos



Compiling OpenMP programs

- Load GNU Compiler module

- ❑ `$ module load gnu7/7.3.0`

- Compiling

- ❑ `$ gcc -o hello-openmp -fopenmp hello-openmp.c`

- Automate compiling using Makefile

- ❑ Edit a Makefile

- ❑ `$ make`



Running OpenMP programs

- Create job submission script, e.g.

```
#!/bin/sh
#$ -V
#$ -cwd
#$ -S /bin/bash
#$ -N OpenMP_Test_Job
#$ -o $JOB_NAME.o$JOB_ID
#$ -e $JOB_NAME.e$JOB_ID
#$ -q omni
#$ -pe sm 36
#$ -l h_vmem=5.3G
#$ -l h_rt=48:00:00
#$ -P quanah
```

```
./hello-openmp
```

- Submit job
 - ❑ **qsub** <job submission script>
- Check job status
 - ❑ Command: **qstat**



Checking Output and Debugging Failed Jobs

- Job output
 - ❑ Standard: `$JOB_NAME.o$JOB_ID`
 - ❑ Error: `$JOB_NAME.e$JOB_ID`

- When debugging:
 - ❑ Check the output files for errors
 - ❑ Check the output of `qacct -j <job_ID>`
 - failed
 - exit_status
 - maxvmem
 - start_time & end_time (<runtime limit)
 - low



Demos

- Source code samples: Please checkout source code samples with executing the following command:

```
git clone  
https://discl.cs.ttu.edu/gitlab/yongchen/cs4379  
cs5379.git
```

- If you have already checked out a copy of the repo earlier, you can run the following command to update to the latest source code repo:

```
git pull
```



More Resources: OpenMP Standard

- Standard developed from 1997, with initial members of Intel, IBM, Compaq, HP, SGI, Sun Microsystems, DOE
- OpenMP website: openmp.org
 - ❑ API specifications, FAQ, presentations, discussions, media releases, calendar, membership application and more...
- OpenMP specifications
 - ❑ <https://www.openmp.org/specifications/>
- OpenMP 5.0 Complete Specifications
 - ❑ <https://www.openmp.org/wp-content/uploads/OpenMP-API-Specification-5.0.pdf>
 - ❑ Useful reference
- Book: “Using OpenMP: Portable Shared Memory Parallel Programming”





Readings

- Reference book ITPC – Chapter 7, 7.10
- OpenMP Programming, by Blaise Barney, Lawrence Livermore National Laboratory: <https://computing.llnl.gov/tutorials/openMP/>
- OpenMP 5.0 Complete Specifications
 - <https://www.openmp.org/wp-content/uploads/OpenMP-API-Specification-5.0.pdf>



Questions?

Questions/Suggestions/Comments are always welcome!

Write me: yong.chen@ttu.edu

Call me: 806-834-0284

See me: ENGCTR 315

If you write me an email for this class, please start the email subject with [CS4379] or [CS5379].