



# **CS4379: Parallel and Concurrent Programming**

## **CS5379: Parallel Processing**

### **Lecture 22**

**Dr. Yong Chen**

**Associate Professor**

**Computer Science Department**

**Texas Tech University**

## Lecture Video

- Please view the lecture video either from Teams or from the below link:

<https://texastechuniversity.sharepoint.com/sites/CS4379-CS5379/Shared%20Documents/General/Lecture22.mp4>

## Course Info

- **Lecture Time:** TR, 12:30-1:50
- **Lecture Location:** ECE 217
- **Sessions:** CS4379-001, CS4379-002, CS5379-001, CS5379-D01
- **Instructor:** Yong Chen, Ph.D., Associate Professor
- **Email:** yong.chen@ttu.edu
- **Phone:** 806-834-0284
- **Office:** Engineering Center 315
- **Office Hours:** 2-4 p.m. on Wed., or by appointment
- **TA:** Mr. Ghazanfar Ali, [Ghazanfar.Ali@ttu.edu](mailto:Ghazanfar.Ali@ttu.edu)
- **TA Office hours:** Tue. and Fri., 2-3 p.m., or by appointment
- **TA Office:** EC 201 A
- **More info:**
  - <http://www.myweb.ttu.edu/yonchen>
  - <http://discl.cs.ttu.edu>; <http://cac.ttu.edu/>; <http://nsfcac.org>

# Outline


- Questions?
- Principles of Message Passing Programming
- MPI (Message Passing Interface) Basics
- Blocking v.s. Non-blocking

# Principles of Message-Passing Programming

- The logical view of a distributed-address-space parallel computer supporting the message-passing paradigm consists of  $p$  processes, each with its own exclusive address space.
- Each data element must belong to some process' address space; hence, data must be explicitly partitioned and placed.
- All interactions (read-only or read and write) require cooperation of two (or more) processes - the process that has the data and the process that wants to access the data.
- Most message-passing programs are written using the single program multiple data (SPMD) model.

# Principles of Message-Passing Programming

- Message passing model is for communication among processes (with separate address spaces)
- Interprocess communication consists of
  - ❑ Movement of data from one process's address space to another's
  - ❑ Synchronization



# What is MPI?

- Message Passing Interface (MPI)
- *A message-passing library specification for parallel computers*
  - Not a language or compiler specification
  - Not a specific implementation or product
- Full-featured
- Designed for
  - Application developers
  - Library writers
  - Tool developers

# Where Did MPI Come From?

- Early vendor systems (Intel's NX, IBM's EUI) were not portable
- Early portable systems (PVM, p4, Chameleon) were mainly research efforts
  - Did not address the full spectrum of message-passing issues
  - Lacked vendor support
  - Were not implemented at the most efficient level
- The MPI Forum organized in 1992 with broad participation by:
  - Vendors: IBM, Intel, TMC, SGI, Convex, Meiko
  - Portability library writers: PVM, p4
  - Users: application scientists and library writers
  - MPI-1 finished in 18 months



# Important Considerations while Using MPI

- All parallelism is explicit: the programmer is responsible for correctly identifying parallelism and implementing parallel algorithms using MPI constructs.

# Hello World (C)

```
#include "mpi.h"
#include <stdio.h>

int main( argc, argv )
int argc;
char *argv[];
{
    int rank, size;
    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &rank );
    MPI_Comm_size( MPI_COMM_WORLD, &size );
    printf( "I am %d of %d\n", rank, size );
    MPI_Finalize();
    return 0;
}
```


# Hello World (Fortran)

```
program main
include 'mpif.h'
integer ierr, rank, size

call MPI_INIT( ierr )
call MPI_COMM_RANK( MPI_COMM_WORLD, rank, ierr )
call MPI_COMM_SIZE( MPI_COMM_WORLD, size, ierr )
print *, 'I am ', rank, ' of ', size
call MPI_FINALIZE( ierr )
end
```

# MPI Basic Send/Receive

- We need to fill in the details in



- Things that need specifying:
  - ❑ How will “data” be described?
  - ❑ How will processes be identified?
  - ❑ How will the receiver recognize/screen messages?
  - ❑ What will it mean for these operations to complete?

## Some Basic Concepts

- Processes can be collected into *groups*.
- Each message is sent in a *context*, and must be received in the same context.
- A group and context together form a *communicator*.
- A process is identified by its *rank* in the group associated with a communicator.
- There is a default communicator whose group contains all initial processes, called **MPI\_COMM\_WORLD**.

# MPI Datatypes

- The data in a message to sent or received is described by a triple (address, count, datatype), where
- An MPI *datatype* is recursively defined as:
  - predefined, corresponding to a data type from the language (e.g., MPI\_INT, MPI\_DOUBLE\_PRECISION)
  - a contiguous array of MPI datatypes
  - a strided block of datatypes
  - an indexed array of blocks of datatypes
  - an arbitrary structure of datatypes
- There are MPI functions to construct custom datatypes, such an array of (int, float) pairs, or a row of a matrix stored columnwise.

# MPI Tags

- Messages are sent with an accompanying user-defined integer *tag*, to assist the receiving process in identifying the message.
- Messages can be screened at the receiving end by specifying a specific tag, or not screened by specifying **MPI\_ANY\_TAG** as the tag in a receive.
- Some non-MPI message-passing systems have called tags “message types”. MPI calls them tags to avoid confusion with datatypes.

## MPI Basic (Blocking) Send

MPI\_SEND (address, count, datatype, dest, tag, comm)

- The message buffer is described by (**address, count, datatype**).
- The target process is specified by **dest**, which is the rank of the target process in the communicator specified by **comm**.
- When this function returns, the data has been delivered to the system and the buffer can be reused. The message may not have been received by the target process.



## MPI Basic (Blocking) Receive

MPI\_RECV(address, count, datatype, source, tag, comm, status)

- Waits until a matching (on **source** and **tag**) message is received from the system, and the buffer can be used.
- **source** is rank in communicator specified by **comm**, or **MPI\_ANY\_SOURCE**.
- **status** contains further information
- receiving fewer than **count** occurrences of **datatype** is OK, but receiving more is an error.

# Status Object

- The status object is used after completion of a receive to find the actual length, source, and tag of a message
- Status object is MPI-defined type and provides information about:
  - The source process for the message (status.source)
  - The message tag (status.tag)
- The number of elements received is given by:

**int MPI\_Get\_count( MPI\_Status \*status, MPI\_Datatype datatype, int \*count )**

**status** return status of receive operation (Status)

**datatype** datatype of each receive buffer element (handle)

**count** number of received elements (integer)(OUT)

# MPI is Simple

- Many parallel programs can be written using just these six functions, only two of which are non-trivial:
  - ❑ `MPI_INIT` - initialize the MPI library (must be the first routine called)
  - ❑ `MPI_COMM_SIZE` - get the size of a communicator
  - ❑ `MPI_COMM_RANK` - get the rank of the calling process in the communicator
  - ❑ `MPI_SEND` - send a message to another process
  - ❑ `MPI_RECV` - send a message to another process
  - ❑ `MPI_FINALIZE` - clean up all MPI state (must be the last MPI function called by a process)
- For performance, however, you need to use other MPI features

# Outline

- Questions?
- Principles of Message Passing Programming
- MPI (Message Passing Interface) Basics
- Blocking v.s. Non-blocking

# Blocking v/s Non-blocking modes

- Blocking mode:
  - Return from the routine implies completion.
- Non-Blocking mode:
  - Routine returns immediately, doesn't imply completion, the completion needs to be tested for
  - Primarily used to overlap computation with communication and exploit possible performance gains
- “Completion” means that memory locations used in the message transfer can be safely accessed for reuse.
  - Safe means that modifications will not affect the data intended for the receive task.
  - For “send” completion implies variable sent can be reused/modified
    - Nothing is said whether the message has been delivered/received
  - For “receive” variable received can be read.

# Blocking Communication

- In Blocking communication
  - MPI\_SEND does not complete until buffer is emptied (available for reuse)
  - MPI\_RECV does not complete until buffer is filled (available for use)
- A process sending data will be blocked until data in the send buffer is emptied
- A process receiving data will be blocked until the receive buffer is filled
- Completion of communication generally depends on the message size and the system buffer size
- Blocking communication is simple to use but can be prone to deadlocks

```
    If (myrank .eq. 0) then
```

```
        Call mpi_send(..)
```

```
        Call mpi_recv(...)
```

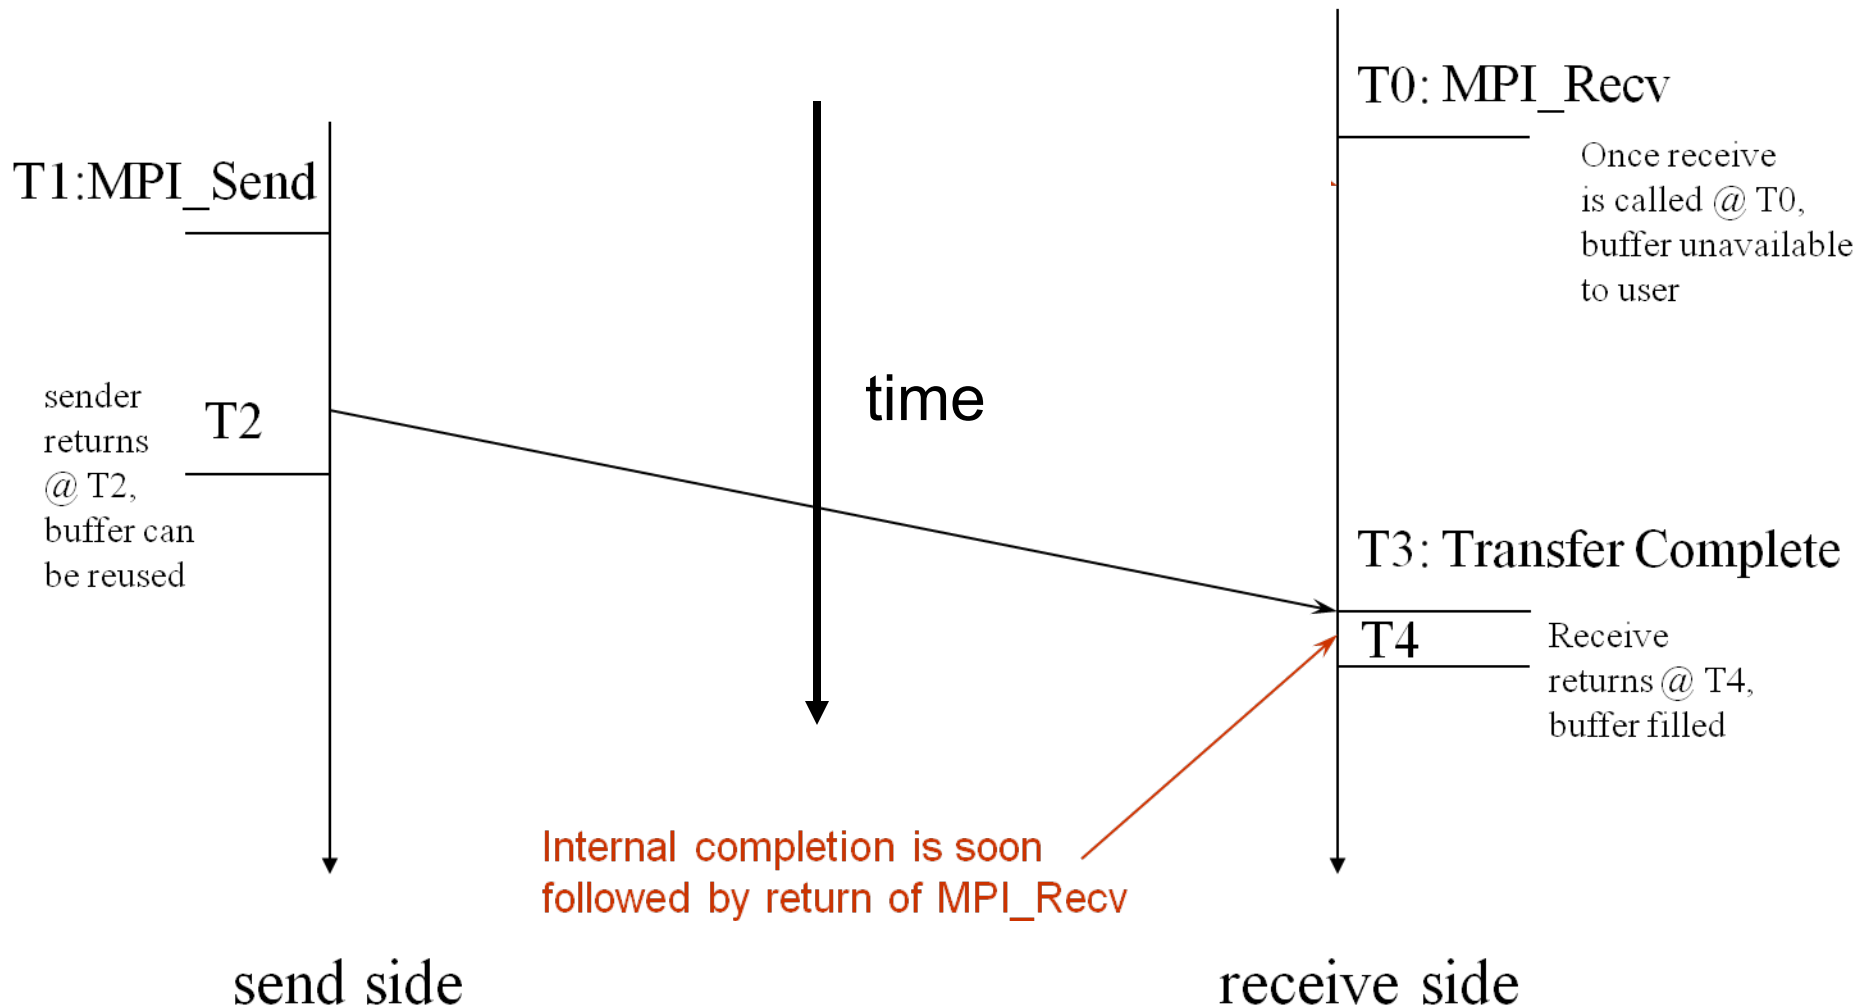
```
Can have deadlocks → Else
```

```
    Call mpi_send(...) ← UNLESS you reverse send/recv
```

```
    Call mpi_recv(...)
```

```
Endif
```

# Blocking Send-Receive Diagram



# Non-Blocking Communication

- Non-blocking operations return (immediately) “request handles” that can be waited on and queried
  - ❑ **MPI\_ISEND**( start, count, datatype, dest, tag, comm, request )
  - ❑ **MPI\_IRECV**( start, count, datatype, src, tag, comm, request )
  - ❑ **MPI\_WAIT**( request, status ) -> Blocking!
- Non-blocking operations allow overlapping computation and communication.
- One can also test without waiting using MPI\_TEST
  - ❑ **MPI\_TEST**( request, flag, status )
- Anywhere you use MPI\_Send or MPI\_Recv, you can use the pair of MPI\_Isend/MPI\_Wait or MPI\_Irecv/MPI\_Wait



# Multiple Completions

- It is sometimes desirable to wait on multiple requests:

```
MPI_Waitall(count, array_of_requests,  
            array_of_statuses)
```


```
MPI_Waitany(count, array_of_requests,  
            &index, &status)
```

```
MPI_Waitsome(count, array_of_requests,  
            array_of_indices, array_of_statuses)
```

- There are corresponding versions of **test** for each of these.

# Non-Blocking Send-Receive Diagram

High Performance Implementations  
Offer Low Overhead for Non-blocking Calls



# Non-Blocking Communication

- Avoid deadlocks
- Post all requests and let matches happen

```
If (myrank .eq. 0) then
    Call mpi_isend(..)
    Call mpi_irecv(...)
Else
    Call mpi_isend(...)
    Call mpi_irecv(....)
Endif
```

# Message Completion and Buffering

- For a communication to succeed:
  - ❑ Sender must specify a valid destination rank
  - ❑ Receiver must specify a valid source rank
  - ❑ The communicator must be the same
  - ❑ Tags must match
  - ❑ Receiver's buffer must be large enough
- A send has completed when the user supplied buffer can be reused

```
*buf = 3;
```

```
MPI_Send (buf, 1, MPI_INT ...)
```

```
*buf = 4; /*OK, receiver will always receive 3*/
```

```
*buf = 3;
```

```
MPI_Isend (buf, 1, MPI_INT ...)
```

```
*buf = 4; /*Not certain if receiver gets 3 or 4*/
```

```
MPI_Wait(...);
```

- Just because the send completes does not mean that the receive has completed
  - ❑ Message may be buffered by the system
  - ❑ Message may still be in transit

# Readings

- Reference book ITPC – Chapter 6, 6.1-6.2



## Questions?

Questions/Suggestions/Comments are always welcome!

Write me: [yong.chen@ttu.edu](mailto:yong.chen@ttu.edu)

Call me: 806-834-0284

See me: ENGCTR 315

*If you write me an email for this class, please start the email subject with [CS4379] or [CS5379].*