



CS4379: Parallel and Concurrent Programming

CS5379: Parallel Processing

Lecture 23

Dr. Yong Chen

Associate Professor

Computer Science Department

Texas Tech University

Lecture Video

- Please view the lecture video either from Teams or from the below link:

<https://texastechuniversity.sharepoint.com/sites/CS4379-CS5379/Shared%20Documents/General/Lecture23.mp4>

Course Info

- **Lecture Time:** TR, 12:30-1:50
- **Lecture Location:** ECE 217
- **Sessions:** CS4379-001, CS4379-002, CS5379-001, CS5379-D01
- **Instructor:** Yong Chen, Ph.D., Associate Professor
- **Email:** yong.chen@ttu.edu
- **Phone:** 806-834-0284
- **Office:** Engineering Center 315
- **Office Hours:** 2-4 p.m. on Wed., or by appointment
- **TA:** Mr. Ghazanfar Ali, Ghazanfar.Ali@ttu.edu
- **TA Office hours:** Tue. and Fri., 2-3 p.m., or by appointment
- **TA Office:** EC 201 A
- **More info:**
 - <http://www.myweb.ttu.edu/yonchen>
 - <http://discl.cs.ttu.edu>; <http://cac.ttu.edu/>; <http://nsfcac.org>

Outline

- Questions?
- Collective Communications in MPI
- More on message passing: communication modes and avoiding deadlocks
- Datatypes in MPI

Collective Communications in MPI

- Collective operations are called by all processes in a communicator
- No message tags used
- In many applications, point-to-point can be replaced by collective communication, improving both simplicity and efficiency
 - Let the internal implementation optimize the communication for you
- Three broad classes:
 - **Synchronization**: barrier
 - **Data movement routines**: broadcast, gather, scatter
 - **Global computation routines**: reduction, scan

Barrier Routine

- Used to **synchronize execution** of a group of processes:

int MPI_Barrier(MPI_Comm comm);

- A barrier is a simple way to separate two phases of computation to ensure that messages in two phases do not interact

Data Movement Routines

- Broadcast routine implements a **one-to-all broadcast** where a single named process (root) sends the same data to all other processes

int MPI_Bcast (void *buffer, int count, MPI_Datatype datatype, int root, MPI_Comm comm)

buffer starting address of buffer (choice)

count number of entries in buffer (integer)

datatype data type of buffer (handle)

root rank of broadcast root (integer)

comm communicator (handle)

Data Movement Routines

*int MPI_Gather (void *sendbuf, int sendcnt, MPI_Datatype sendtype, void *recvbuf,
int recvcnt, MPI_Datatype recvtype, int root, MPI_Comm comm)*

*int MPI_Scatter (void *sendbuf, int sendcnt, MPI_Datatype sendtype, void *recvbuf,
int recvcnt, MPI_Datatype recvtype, int root, MPI_Comm comm)*

sendbuf starting address of send buffer (choice)

sendcount number of elements in send buffer (integer)

sendtype data type of send buffer elements (handle)

recvcnt number of elements for any single receive (integer, significant only at root)


recvtype data type of recv buffer elements (significant only at root) (handle)

root rank of receiving/sending process (integer)

comm communicator (handle)

recvbuf address of receive buffer (choice, significant only at root) (OUT)

Illustration of MPI Communication Functions




Collective Communication Operations

- The all-to-all personalized communication operation is performed by:

```
int MPI_Alltoall(void *sendbuf, int sendcount,  
    MPI_Datatype senddatatype, void *recvbuf,  
    int recvcount, MPI_Datatype recvdatatype,  
    MPI_Comm comm)
```

- Each process sends to process i sendcount contiguous elements of type senddatatype starting from the $i * \text{sendcount}$ location of its sendbuf
- The data that are received are stored in the recvbuf array
- Each process receives from process i recvcount elements of type recvdatatype and stores them in its recvbuf array starting at location $i * \text{recvcount}$

All-to-all Personalized Communication



Reduction Operations

- Reduction operations **combine the values** provided in the input buffer of each process **using a specified operation OP**, and return combined value into output buffer of single root process or output buffer of all processes

int MPI_Reduce (void *sendbuf, void *recvbuf, int count, MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm comm)

int MPI_Allreduce (void *sendbuf, void *recvbuf, int count, MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)

sendbuf address of send buffer (choice)

count number of elements in send buffer (integer)

datatype data type of elements of send buffer (handle)


op reduce operation (handle)

root rank of root process (integer)


comm communicator (handle)

recvbuf address of receive buffer (choice, significant only at root) (OUT)


Illustration of Reduction Operations




MPI_REDUCE with MPI_MIN, root = 0 :



MPI_ALLREDUCE with MPI_MIN:



MPI_REDUCE with MPI_SUM, root = 1 :



MPI Built-in Collective Computation Operations

■ MPI_Max	Maximum
■ MPI_Min	Minimum
■ MPI_Prod	Product
■ MPI_Sum	Sum
■ MPI_Land	Logical and
■ MPI_Lor	Logical or
■ MPI_Lxor	Logical exclusive or
■ MPI_Band	Binary and
■ MPI_Bor	Binary or
■ MPI_Bxor	Binary exclusive or
■ MPI_Maxloc	Maximum and location
■ MPI_Minloc	Minimum and location

Collective Communication Operations

- The operation `MPI_MAXLOC` combines pairs of values (v_i, l_i) and returns the pair (v, l) such that v is the maximum among all v_i 's and l is the corresponding l_i (if there are more than one, it is the smallest among all these l_i 's).
- `MPI_MINLOC` does the same, except for minimum value of v_i .

Value	15	17	11	12	17	11
Process	0	1	2	3	4	5

`MinLoc(Value, Process) = (11, 2)`

`MaxLoc(Value, Process) = (17, 1)`


An example use of the `MPI_MINLOC` and `MPI_MAXLOC` operators.

Example: Calculating Pi

- One way to calculate Pi:

$$\pi = \int_0^1 \frac{4}{1+x^2} dx$$

- Calculating Pi via numerical integration
 - ▣ Divide and assign to processes
 - ▣ Each process calculates partial sum
 - ▣ Add all the partial sums together to get Pi



Example: PI in C (1/2)

```
#include "mpi.h"
#include <math.h>
int main(int argc, char *argv[])
{
    int done = 0, n, myid, numprocs, i, rc;
    double PI25DT = 3.141592653589793238462643;
    double mypi, pi, width, sum, x, a;
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);
    while (!done) {
        if (myid == 0) {
            printf("Enter the number of intervals: (0 quits) ");
            scanf("%d", &n);
        }

        MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);
        if (n == 0) break;
    }
}
```

message/data
broadcasted
root process

Example: PI in C (2/2)

```
width = 1.0 / (double) n;
```

```
sum = 0.0;
```

```
for (i = myid + 1; i <= n; i += numprocs) {
```

```
    x = width * ((double)i - 0.5);
```

```
    sum += 4.0 / (1.0 + x*x);
```

```
}
```

```
myypi = width * sum;
```

```
MPI_Reduce(&myypi, &pi, 1, MPI_DOUBLE, MPI_SUM, 0,  
           MPI_COMM_WORLD);
```

```
if (myid == 0)
```

```
    printf("pi is approximately %.16f, Error is %.16f\n",  
           pi, fabs(pi - PI25DT));
```

```
}
```

```
MPI_Finalize();
```

```
return 0;
```

```
}
```

send buffer
receive buffer
operation
root process

More on MPI Collective Routines

- Many Routines: **Allgather, Allgatherv, Allreduce, Alltoall, Alltoallv, Bcast, Gather, Gatherv, Reduce, ReduceScatter, Scan, Scatter, Scatterv**
- **All** versions deliver results to all participating processes.
- **V** versions allow the chunks to have different sizes.
- **Allreduce, Reduce, ReduceScatter, and Scan** take both built-in and **user-defined combiner/reduction functions**.
 - Create your own collective computations with:
MPI_Op_create(user_fcn, commutes, &op);
MPI_Op_free(&op);

Outline

- Questions?
- Collective Communications in MPI
- More on message passing: communication modes and avoiding deadlocks
- Datatypes in MPI

Communication Modes

- MPI has multiple *communication modes* for sending messages:
 - **Synchronous mode (MPI_Ssend)**: the send waits for a matching receive has begun. (Unsafe programs deadlock.)
 - **Buffered mode (MPI_Bsend)**: the user supplies a buffer to the system for its use. (User allocates enough memory to make an unsafe program safe.)
 - **Ready mode (MPI_Rsend)**: user guarantees that a matching receive has been posted.
 - Allows access to fast protocols
 - Undefined behavior if matching receive not posted
- Non-blocking versions (MPI_Issend, etc.)
- MPI_Recv/MPI_IRecv receives messages sent in any mode

Flavors of Communication

- For a send operation there are:
 - ❑ 4 communication modes: standard, ready, synchronous, buffered
 - ❑ 2 blocking modes: blocking, nonblocking
 - ❑ $4 * 2 = 8$ types of sends

- For a receive operation there are:
 - ❑ 1 communication mode: standard
 - ❑ 2 blocking modes: blocking, nonblocking
 - ❑ $1 * 2 = 2$ types of receive

Naming Conventions

- Send routines:

Comm Mode	Blocking	Non-blocking
Standard	MPI_Send	MPI_Isend
Ready	MPI_Rsend	MPI_Irsend
Synchronous	MPI_Ssend	MPI_Issend
Buffered	MPI_Bsend	MPI_Ibsend

- Receive routines:

Comm Mode	Blocking	Non-blocking
Standard	MPI_Recv	MPI_Irecv

- Any type of receive routine can be used to receive messages from any type of send routine

Send/Receive Operations

- In many applications, processes send to one process while receiving from another
- Deadlock may arise if care is not taken
- MPI provides routines for such send/receive operations

int MPI_Sendrecv(void *sendbuf, int sendcount, MPI_Datatype sendtype, int dest, int sendtag, void *recvbuf, int recvcount, MPI_Datatype recvtype, int source, int recvtag, MPI_Comm comm, MPI_Status *status)

Deadlocks

- What happens with this code?

Process 0

Process 1

Send (1)

Send (0)

Recv (1)

Recv (0)

- MPI_Send and MPI_Recv are blocking comm:
 - ❑ MPI_Send does not complete until sending buffer is empty
 - ❑ MPI_Recv does not complete until receiving buffer is full
- Simple, but maybe “unsafe” because completion depends on the message size and availability of system buffers

Solutions to the Unsafe Problem

- Order the operations more carefully:

Process 0	Process 1
Send (1)	Recv (0)
Recv (1)	Send (0)

- Supply receive buffer at same time as send:

Process 0	Process 1
Sendrecv (1)	Sendrecv (0)

More Solutions to the Unsafe Problem

- Supply own space as buffer for send

Process 0	Process 1
Bsend(1)	Bsend(0)
Recv(1)	Recv(0)

- Use non-blocking operations:

Process 0	Process 1
Isend(1)	Isend(0)
Irecv(1)	Irecv(0)
Waitall	Waitall

Outline

- Questions?
- Collective Communications in MPI
- More on message passing: communication modes and avoiding deadlocks
- Datatypes in MPI

Datatypes in MPI

- MPI datatypes have two purposes:
 - Heterogeneity
 - Noncontiguous data

- Basic vs. derived datatype:
 - Basic datatype
 - Derived datatype
 - Contiguous
 - Vector
 - Indexed
 - Hindexed
 - Structure

Basic Datatype in C

MPI Datatype	C Datatype
MPI_BYTE	
MPI_CHAR	signed char
MPI_DOUBLE	double
MPI_FLOAT	float
MPI_INT	int
MPI_LONG	long
MPI_LONG_DOUBLE	long double
MPI_PACKED	
MPI_SHORT	short
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long
MPI_UNSIGNED_SHORT	unsigned short


Additional datatypes defined in MPI 2.2 corresponding to C99 language types `int32_t`, `int64_t`, etc.

Typemaps in MPI

- In MPI, a datatype is represented as a **typemap**

$$typemap = (type_0, disp_0), \dots, (type_{n-1}, disp_{n-1})$$

- **Extent** of a datatype



- An artificial extent can be set by using MPI_UB and MPI_LB

Typemaps in MPI (cont.)

- Example:
 - ❑ (int,0),(char,4) is a typemap
 - ❑ The extent of this typemap is 5

CONTIGUOUS Datatype

MPI_Type_contiguous(count, oldtype, &newtype)

MPI_Type_commit(&newtype)

- Assume an original datatype oldtype has typemap (double,0), (char,8), then

MPI_Type_contiguous(3,oldtype,&newtype);

- To actually send such a data use the sequence of calls:

MPI_Type_contiguous(count,datatype,&newtype);

MPI_Type_commit(&newtype);

MPI_Send(buffer,1,newtype,dest,tag,comm);

MPI_Type_free(&newtype);

VECTOR Datatype

MPI_Type_vector(count, blocklength, stride, oldtype, &newtype)

MPI_Type_commit(&newtype)

29	30	31	32	33	34	35
22	23	24	25	26	27	28
15	16	17	18	19	20	21
8	9	10	11	12	13	14
1	2	3	4	5	6	7

MPI_Type_Vector(5,1,7,MPI_DOUBLE,newtype);

MPI_Type_commit(&newtype);

MPI_Send(buffer,1,newtype,dest,tag,comm);


MPI_Type_free(&newtype);

INDEXED Datatype

MPI_Type_indexed(count, &array_of_blocklengths, &array_of_displacements, oldtype, &newtype)


MPI_Type_commit(&newtype)

MPI_Type_indexed(count, blocklengths, displacements, MPI_FLOAT, &indextype);



Hindexed:
MPI_Type_hindexed()
is the same except that
offsets array is
specified in bytes

MPI_Send(&a, 1, indextype, dest, tag, comm);



1 element of indextype

Structure Datatype


*MPI_Type_struct(count, &array_of_blocklengths, &array_of_displacements,
oldtypes, &newtype)*

MPI_Type_commit(&newtype)

```
MPI_Type_struct(count, blocklengths, offsets, oldtypes, &particletype);
```

```
typedef struct { float f1, f2, f3, f4; int n1, n2;} Particle;
```

```
Particle particles[NELEM];
```



```
MPI_Send(particles, NELEM, particletype, dest, tag, comm);
```

6.0	7.0	8.0	9.0	10	11
-----	-----	-----	-----	----	----

1 element of particletype

Readings

- Reference book ITPC – Chapter 6, 6.3-6.6



Questions?

Questions/Suggestions/Comments are always welcome!

Write me: yong.chen@ttu.edu

Call me: 806-834-0284

See me: ENGCTR 315

If you write me an email for this class, please start the email subject with [CS4379] or [CS5379].